




Auto-Vectorization in Compilers: Leveraging SIMD for High- Performance Computing





Content



01

Background


02

Predicated SSA

[PLDI22](#)

| Compilation and Compiler

- **Compilation** is the process of translating **high-level programming language** (e.g. C++, Python) source code into **low-level machine-executable** (e.g., assembly).
- A **compiler** is a software tool that translates source code into target machine code, often applying optimizations to improve performance and efficiency.



| Some Concepts in Compiler

Control Flow Analysis


- **Basic Block** is a sequence of consecutive instructions in a program. These instructions execute in sequence under same condition.
 - Basic blocks form a program's Control Flow Graph (CFG), where edges represent possible execution paths between blocks.

Dataflow Analysis

- **Static Single Assignment (SSA)** is an intermediate representation (IR) used in modern compilers to simplify optimization and analysis.
 - **LLVM IR follows SSA form.**
 - Key Property: Each variable is assigned (defined) exactly once in the program.
 - *Phi (Φ) functions* handle values coming from multiple control paths.

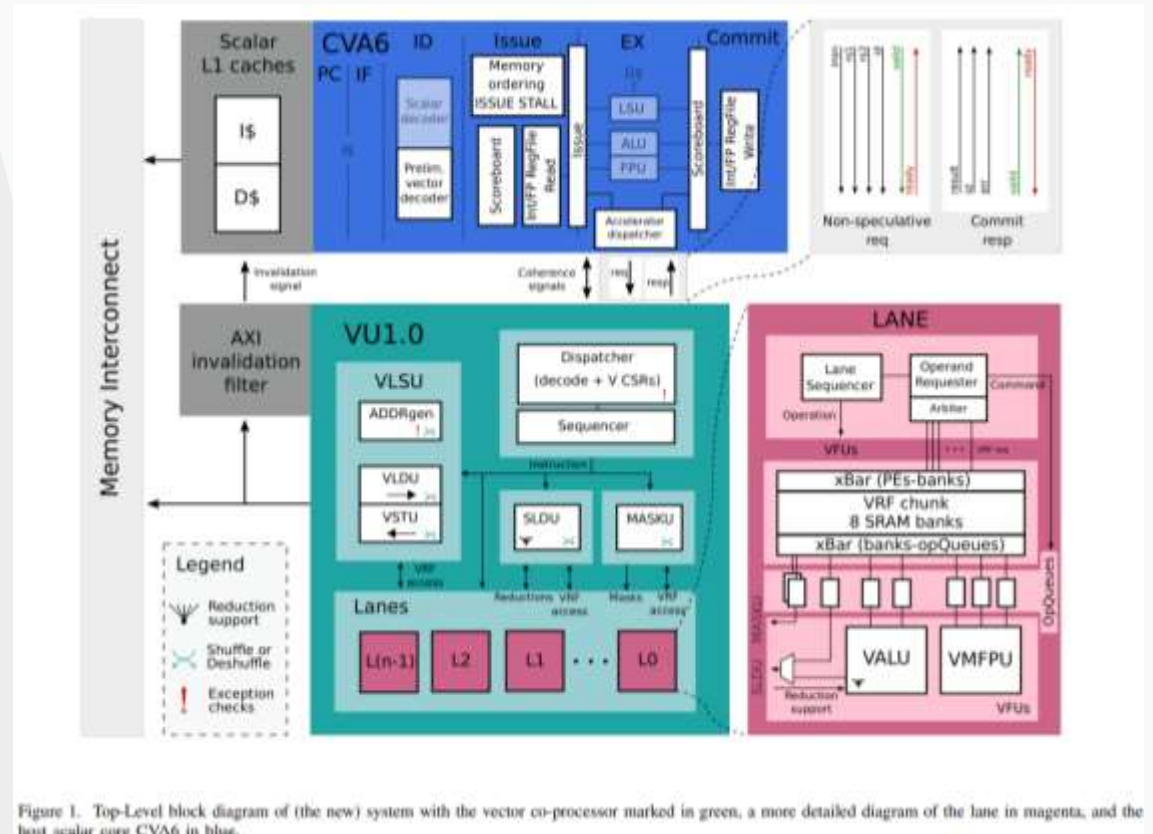
```
x = 5;  
if (condition) {  
    x = x + 1;  
} else {  
    x = x - 1;  
}  
y = x * 2;
```

```
x1 = 5;  
if (condition) {  
    x2 = x1 + 1;  
} else {  
    x3 = x1 - 1;  
}  
x4 =  $\phi$ (x2, x3);  
y1 = x4 * 2;
```



| SIMD

- processors now include SIMD registers and instructions, to enable more parallelism.
- Single Instruction Multiple Data
- vector registers are extra wide
 - a 512-bit wide register can hold 16 32-bit words.
- Instructions on these registers perform the **same operation** on each of the packed data types.




M. Perotti, M. Cavalcante, N. Wistoff, R. Andri, L. Cavigelli and L. Benini, "A "New Ara" for Vector Computing: An Open Source Highly Efficient RISC-V V 1.0 Vector Processor Design," 2022 IEEE 33rd International Conference on Application-specific Systems, Architectures and Processors (ASAP), Gothenburg, Sweden, 2022, pp. 43-51, doi: 10.1109/ASAP54787.2022.00017. keywords: {Program processors;Source coding;Systems architecture;Computer architecture;Throughput;Vector processors;Kernel;RISC-V;ISA;Vector;Efficiency},

| SIMD

```
void vec_add_rvv(vint32m1_t *a, vint32m1_t *b, vint32m1_t *c, size_t vl) {  
    *c = vadd_vv_i32m1(*a, *b, vl);  
}
```

1020:	0287ec87	vl1re32.v	v25,(a5)
1024:	fe043783	ld	a5,-32(s0)
1028:	0287ec07	vl1re32.v	v24,(a5)
102c:	fd043783	ld	a5,-48(s0)
1030:	0507f057	vsetvli	zero,a5,e32,m1,ta,mu
1034:	039c0c57	vadd.vv	v24,v25,v24
1038:	fd843783	ld	a5,-40(s0)
103c:	02878c27	vs1r.v	v24,(a5)



| Automatic Vectorization

- Loop Vectorize

1024 scalar addition

```
for (i = 0; i < 1024; i++)  
    c[i] = a[i] * b[i];
```



256 vector addition

```
for (i = 0; i < 1024; i += 4)  
    c[i:i+3] = a[i:i+3] * b[i:i+3];
```

- SLP(superword level parallelism) Vectorize

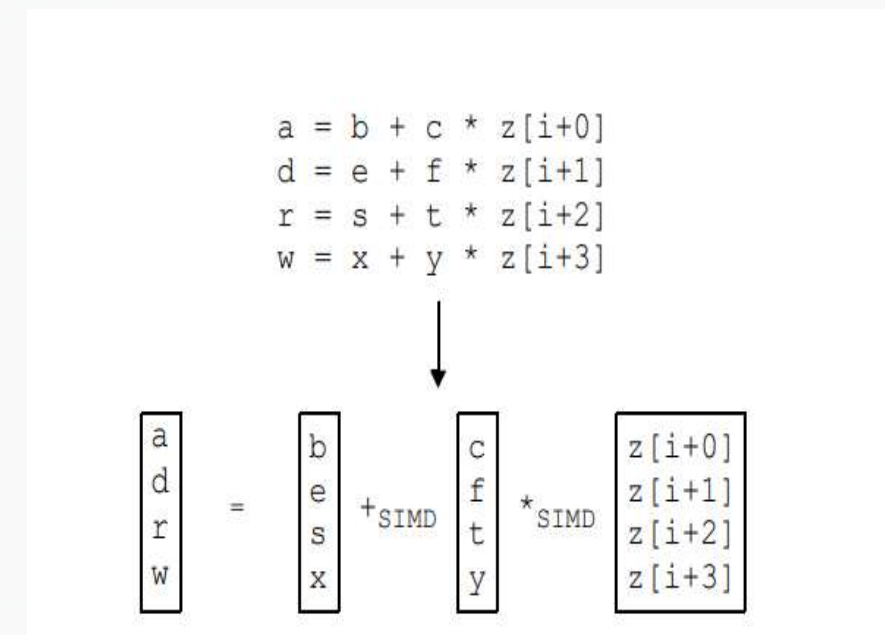
```
a = b + c * z[i+0]  
d = e + f * z[i+1]  
r = s + t * z[i+2]  
w = x + y * z[i+3]
```



```
a  
d  
r  
w  
=  
b  
e  
s  
x  
+SIMD  
c  
f  
t  
y  
*SIMD  
z[i+0]  
z[i+1]  
z[i+2]  
z[i+3]
```


| Auto-Vectorization: SLP

- The operations execute in lock-step.
- The inputs and outputs of the operations reside in **packed storage** (usually implemented as vector registers).
- The operations are **isomorphic**.
- The operations are applied **elementwise**.



All You Need Is Superword-Level Parallelism: Systematic Control-Flow Vectorization with SLP



| All You Need is SLP

- SLP vectorization is simpler: not requiring loop dependence analysis to reason with loop-carried dependence
- SLP vectorization is more flexible: more heuristic regarding new SIMD instructions can be used to extend the search space
- Loop Level Parallelism (LLP) can be exploited with Superword Level Parallelism (SLP) after unrolling with appropriate factor.

```
for (i = 0; i < n; ++i) {  
    a[i] = b[i] + c[i];  
}
```



loop unrolling




```
for (i = 0; i < n; i+=4) {  
    /*  
    After unrolling,  
    before SLP vectorization:  
    a[i]   = b[i] + c[i];  
    a[i+1] = b[i+1] + c[i+1];  
    a[i+2] = b[i+2] + c[i+2];  
    a[i+3] = b[i+3] + c[i+3];  
  
    After SLP vectorization:  
    */  
    a[i:i+4] = b[i:i+4]  
              + c[i:i+4];  
}
```

But SLP's capacity is limited on CFG-based IRs

SLP's capacity is limited on CFG-based IRs

- **SLP vectorization is always limited in one Basic Block**
 - Packing multiple instructions into one is a lot about code motion.
 - Vectorizing instructions in different basic block requires code motion across *block boundaries* and thus requires control flow transformation.

```
for (int i = 0; i < n; i++)  
  if (arr[i] == keys[0]) {  
    idxs[0] = i;  
    break;  
  }  
  
for (int i = 0; i < n; i++)  
  if (arr[i] == keys[1]) {  
    idxs[1] = i;  
    break;  
  }
```



Need global code motion and CFG reconstruction to exploit the parallelism

| Predicated SSA

An IR to Facilitate Code Motion.

- Transform CFG to straight-line code.
- Encode control flow with **Predicate**, bool expression indicating whether the instruction can execute.

control predicate *Item: Instruction | Loop*

- *Item: Instruction | Loop*
- *[CP]: control predicate*
 - *True* indicate the *Item* executes unconditionally

Loops are defined in a nested manner

- *Mu* instructions and the *latch predicate* preserve the control relation of loop

$fn ::= item_1 : p_1, \dots, item_n : p_n$

$p ::= \text{control predicate}$

$item ::= \text{instruction} \mid loop$

$loop ::= \text{with } v_1 = \mu_1, \dots, v_m = \mu_m \text{ do}$
 $item_1 : p_1, \dots, item_n : p_n$
 while p_{cont}

$\mu ::= \text{mu}(v_{init}, v_{rec})$

$\phi ::= \text{phi}(v_1 : p_1, \dots, v_n : p_n)$

$v ::= \mu \mid \text{instructions} \mid \text{constant} \mid \text{argument}$

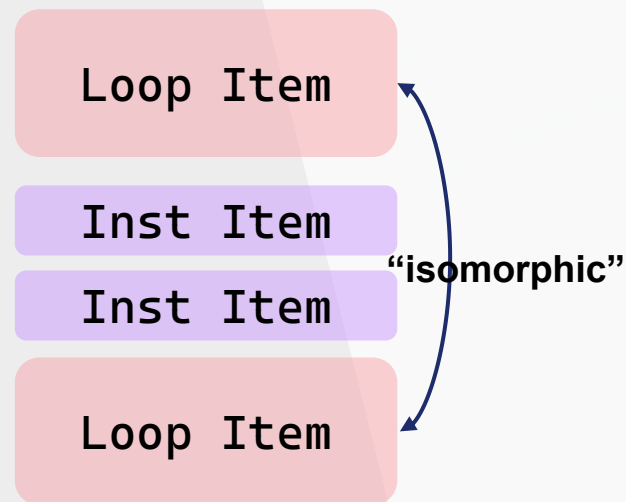

$c ::= \text{branch-conditions}$ (*IR values used for branching*)

$p ::= \text{true} \mid c \mid \bar{c} \mid p_1 \wedge p_2 \mid p_1 \vee p_2$

| Predicated SSA

- Convert CFG of basic block into ‘super BB’ of *Items*.
 - Wider space to do code motion
- Cover control flow information in *Predicate*.
 - **Make vectorization between different loops practical.**
 - Loop Fusion
 - Loop Co-iteration: loops with similar iterate mode can be transferred in to one loop

```
for (int i = 0; i < n; i++)  
  if (arr[i] == keys[0]) {  
    idxs[0] = i;  
    break;  
  }  
  
other code  
  
for (int i = 0; i < n; i++)  
  if (arr[i] == keys[1]) {  
    idxs[1] = i;  
    break;  
  }
```



```
for (int i = 0; i < n; i++)  
  if (arr[i] == keys[0]) {  
    idxs[0] = i;  
    break;  
  }
```

```
for (int i = 0; i < n; i++)  
  if (arr[i] == keys[1]) {  
    idxs[1] = i;  
    break;  
  }
```

```
with i = mu(0, i') do  
  x = arr[i] : true  
  k = keys[0] : true  
  found = cmp eq t, k : true  
  i' = add i, 1 : true  
  lt_n = cmp lt i', n : true  
while not found and lt_n : true  
  idxs[0] = i : found  
with i2 = mu(0, i2') do  
  x2 = arr[i2] : true  
  k2 = keys[1] : true  
  found2 = cmp eq t2, k2 : true  
  i2' = add i2, 1 : true  
  lt_n2 = cmp lt i2', n : true  
while not found2 and lt_n2 : true  
  idxs[1] = i2 : found2
```



Vectorization Pipeline with Predicated SSA

LLVM IR

Convert to PSSA

```
with i = mu(0, i') do
  x = arr[i]           : true
  k = keys[0]          : true
  found = cmp eq t, k  : true
  i' = add i, 1         : true
  lt_n = cmp lt i', n  : true
  while not found and lt_n : true
    idxs[0] = i         : found
  with i2 = mu(0, i2') do
    x2 = arr[i2]        : true
    k2 = keys[1]        : true
    found2 = cmp eq t2, k2 : true
    i2' = add i2, 1      : true
    lt_n2 = cmp lt i2', n : true
    while not found2 and lt_n2 : true
      idxs[1] = i2      : found2
```

Scheduling
(code motion)

```
with i = mu(0, i')
  x = arr[i]           : active1
  x2 = arr[i2]          : active2
  k = keys[0]           : active1
  k2 = keys[1]          : active2
  found = cmp eq t, k    : active1
  found2 = cmp eq t2, k2 : active2
  i' = add i, 1          : active1
  i2' = add i2, 1        : active2

  while active1 or active2 : true
    idxs[0] = i           : found_out'
    idxs[1] = i2          : found_out2'
```

Vector
Instruction
Generation

```
with i = mu(0, i')
  x = arr[i:i+2]        : true
  k = keys[0:2]          : true
  found = vcmp eq x, k   : true
  i' = add i, 1          : active1
  i2' = add i2, 1        : active2


  while active1 or active2 : true
    masked-vstore idxs, i_out', found_out': true
```

Convert to
CFG-based IR

LLVM IR



Evaluation



Evaluation

benchmark for polyhedral optimizations




Figure 14. Speedup over LLVM scalar optimizations (-O3) on PolyBench





Figure 15. Speedup over full LLVM optimizations (-O3 + vectorization) on PolyBench using Polly

| Some Thoughts

- **Vectorization has a lot to do with code motion (scheduling)**
 - Independent instructions in a basic block can be reordered and transform into vector instruction
 - Instructions in different basic block can be viewed as have some kind of dependency.



We care about dependence analysis!

- **Dependence**
 - Control dependence → Contains pseudo dependence
 - Control equivalent instructions in different blocks has pseudo dependence
 - Use-define data dependence → SSA make this trivial globally
 - Memory access data dependence → Gotten from the control flow and the memory access pattern of instructions in each block

| Some Thoughts

- **Dependence**
 - Control dependence —————→ Contains pseudo dependence
 - Use-define data dependence
 - Memory access data dependence
- Control equivalent instructions in different blocks has pseudo dependence

Control dependence can be transferred to data dependence using predication

- That's what PSSA do!
- And emitting the pseudo dependence in loop conditions is the reason why it can perform loop fusion and co-iteration well.
- PSSA is an IR to Facilitate Code Motion.
- Personally, I prefer to call it an IR to facilitate dependence analysis. 🤖



**Thank you
for listening**

