

LLM Acceleration

2023.4.6

Outline

- Challenges
- Inference
- Fine-tuning

Challenges


- Large memory footprint
 - Load model weights to GPU
 - KV Cache
- Long inference Latency

Inference Acceleration

- Quantization
- Pruning
- Off-loading
- Distillation


Quantization

- Converting the weights and activations from fp32 to low bit width
- Post-Training Quantization (PTQ)
- Quantization-Aware Training (QAT)




Quantization

$$\bar{\mathbf{X}}^{\text{INT8}} = \left\lceil \frac{\mathbf{X}^{\text{FP16}}}{\Delta} \right\rceil, \quad \Delta = \frac{\max(|\mathbf{X}|)}{2^{N-1} - 1},$$




(a) per-tensor quantization



(b) per-token + per-channel quantization

Challenges of Quantization on LLM


- The large number of parameters: we prefer PTQ on LLM
- Large outliers: significant degradation of quantization resolution



Configuration	CoLA	SST-2	MRPC	STS-B	QQP	MNLI	QNLI	RTE	GLUE
FP32	57.27	93.12	88.36	89.09	89.72	84.91	91.58	70.40	83.06
W8A8	54.74	92.55	88.53	81.02	83.81	50.31	52.32	64.98	71.03
W32A8	56.70	92.43	86.98	82.87	84.70	52.80	52.44	53.07	70.25
W8A32	58.63	92.55	88.74	89.05	89.72	84.58	91.43	71.12	83.23

Features of Open Source GPT-like LLM

- OPT-175B
with significant outliers, inadequate training, significant redundancy
- BLOOM-176B
less outliers, similar performance with OPT-175B

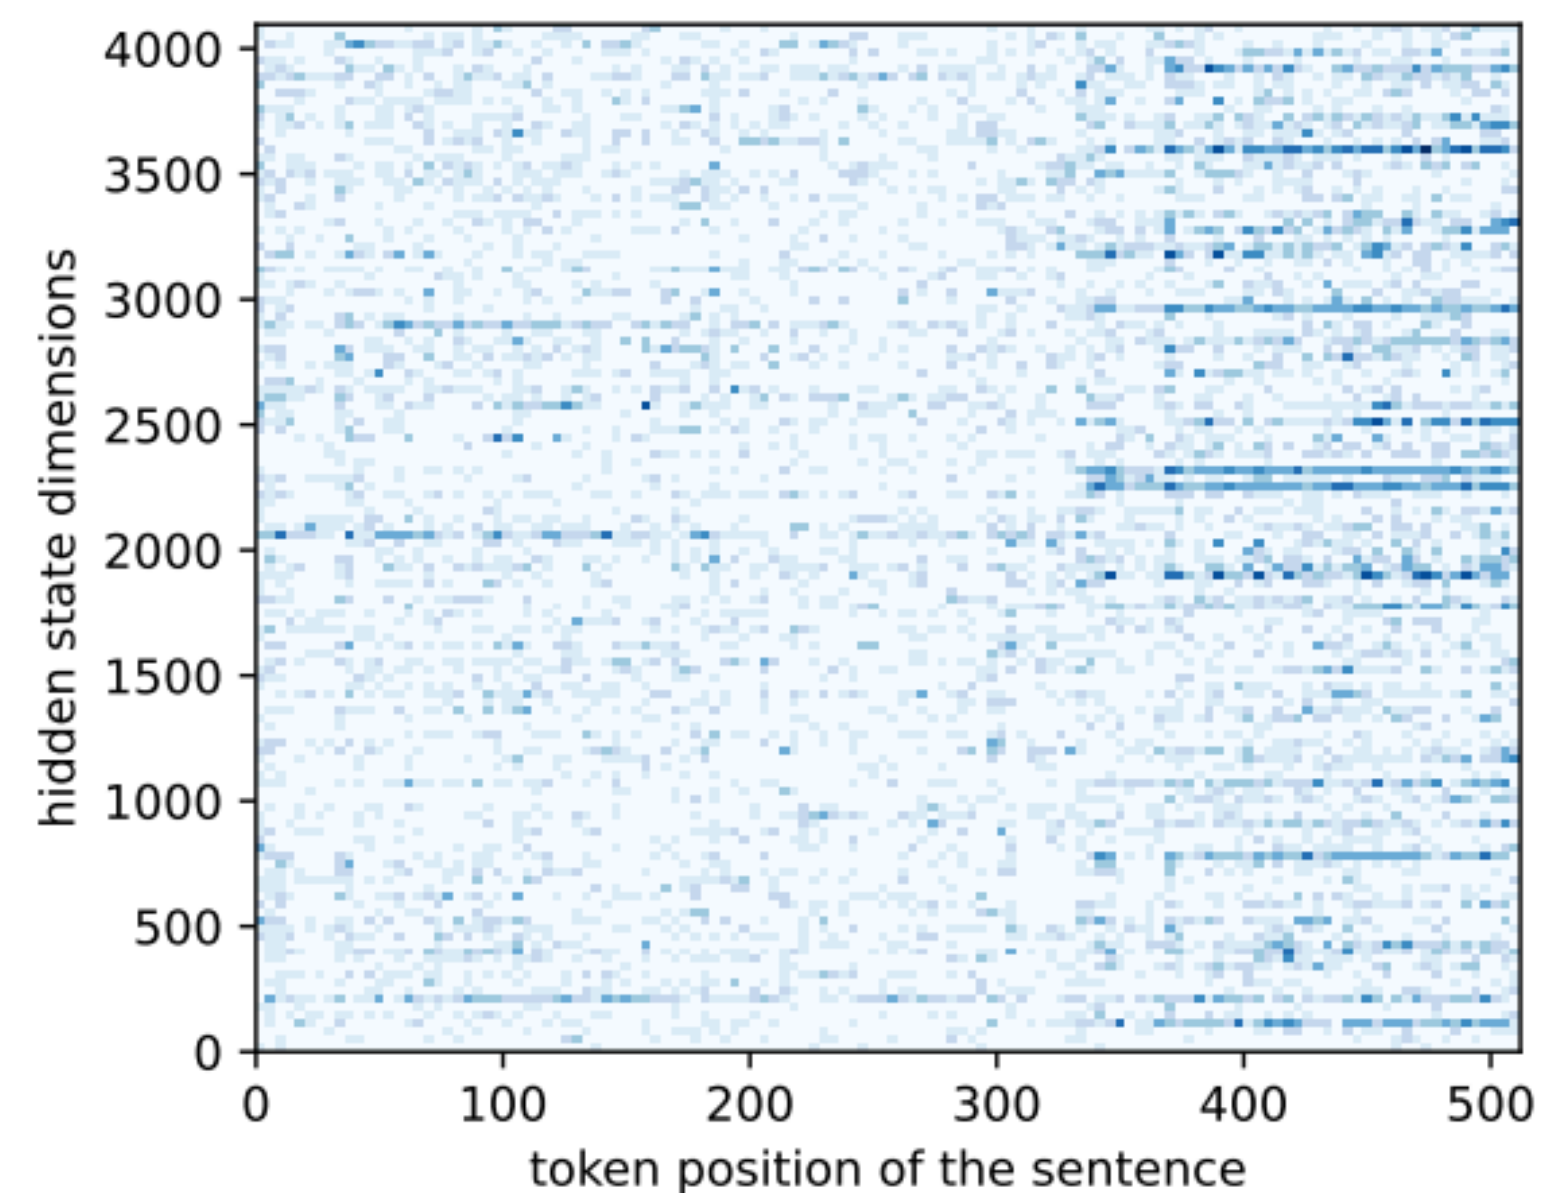
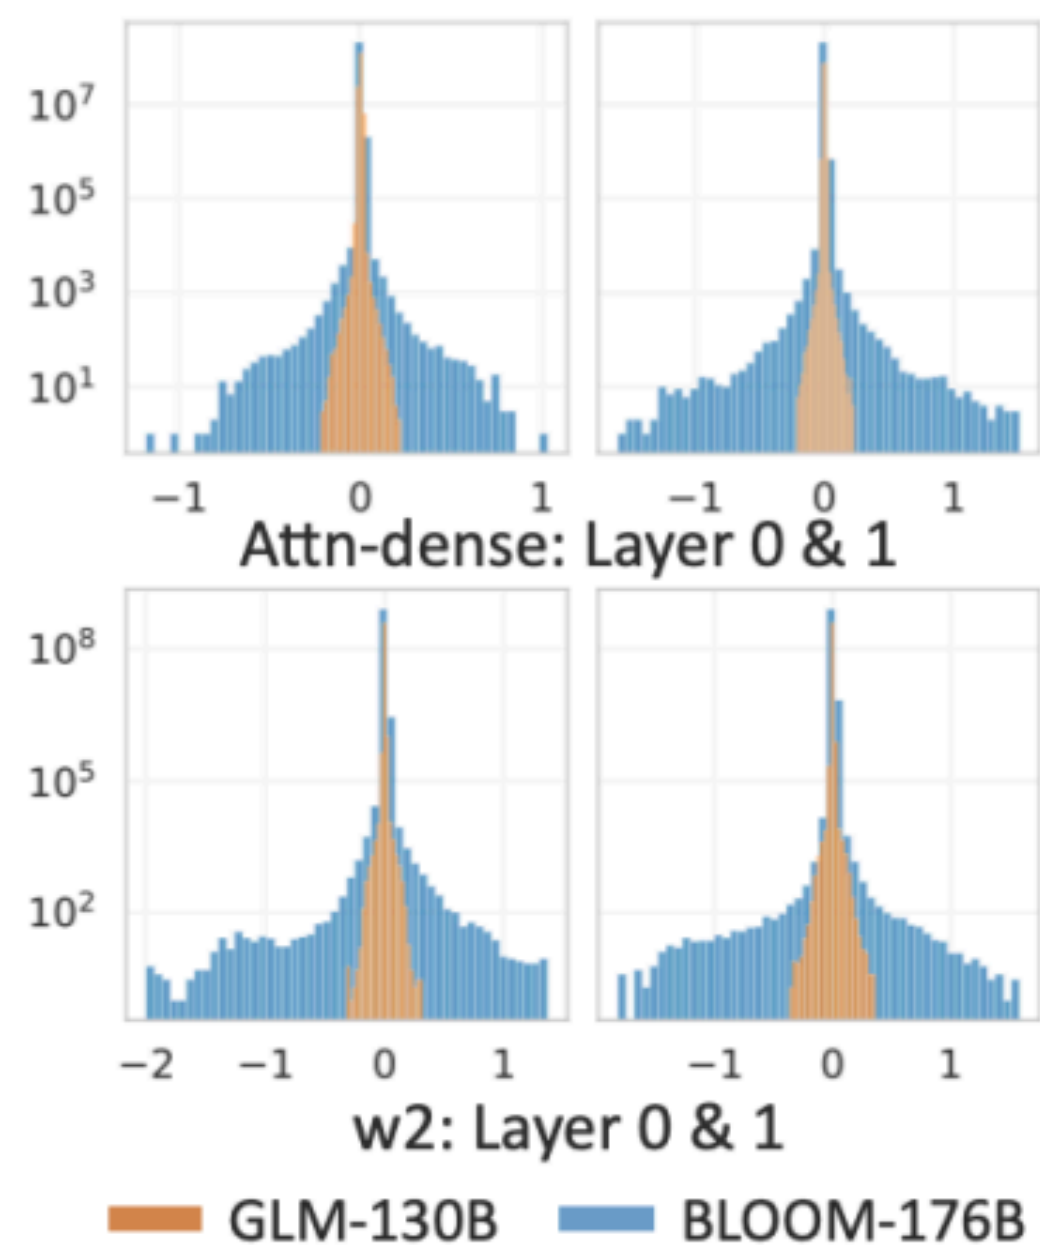


Features of Open Source GPT-like LLM

- GLM-130B

Weight: INT4

Activations: 30% are outliers



Features of Open Source GPT-like LLM

- LLaMA-65B / Alpaca-65B

well-trained, little outliers in both weights and activations

```
layers.10.attention.wq torch.S
tensor([[ 7.8828, -0.0096],
        [ 1.4023,  0.1812],
        [ 0.5474,  0.1475],
        [ 0.0201,  0.1840],
        [-0.2106,  0.2035],
        [-0.0312,  0.1938],
        [ 0.0599,  0.0648],
        [ 0.2041,  0.0465],
        [ 0.1179,  0.1434],
```

```
model.decoder.layers.21.self_attn.k_
tensor([[ 61.5625,  0.0128],
        [ 44.0000,  0.6533],
        [ 24.0625,  0.1434],
        [ 22.3750,  1.4014],
        [ 33.4375,  0.1888],
        [ 36.7188, -0.8003],
        [ 38.3125,  0.4158],
        [ 42.5000,  0.2886],
        [ 29.8438, -0.3481],
```

Types of Quantization

- Quantize both weight and activation
 - Advantages: **accelerate inference** and reduce memory cost
 - Disadvantages: **activation is difficult to maintain precision in low bit width**
 - Related works:
 - LLM.int8(): 99.9% W8A8 LLM loseless
 - Outlier Suppression: W6A6 BERT loseless
 - SmoothQuant: W8A8 LLM loseless



Types of Quantization

- Weight-only quantization
 - Advantages: Lower the memory requirements (People currently prefer)
 - Disadvantages: Fairly slow inference speed
- Related works:
 - OBQ
 - GPTQ (OPT-175B 3bits loseless)
 - llama.cpp (LLaMA-65B on Mbp, LLaMA-7B on Raspberry Pi 4G)
- Is it possible to make hardware support mixed precision multiplication?

LLM.int8() (NIPS 2022)

LLM.int8(): 8-bit Matrix Multiplication for Transformers at Scale


- 1. As the **model size grows to billions of parameters, outliers start to emerge in all transformer layers**, causing failure of simple low-bit quantization.
- 2. **Outliers persist in fixed channels**



LLM.int8() (NIPS 2022)

LLM.int8(): 8-bit Matrix Multiplication for Transformers at Scale

- Isolate the outlier feature dimensions into a 16-bit matrix multiplication while still more than **99.9% of values are multiplied in 8-bit**











LLM.int8() (NIPS 2022)

LLM.int8(): 8-bit Matrix Multiplication for Transformers at Scale

- Perform inference in INT8 LLMs with up to 175B parameters **without any performance degradation**
- **Only the 13B and 175B models have speedups**

SmoothQuant (ICML 2023)


SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models

	LLM (100B+) Accuracy	Hardware Efficiency
ZeroQuant		
Outlier Suppression		
LLM.int8()		
SmoothQuant		


SmoothQuant (ICML 2023)

SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models

- 1. Activations are harder to quantize than weights
- 2. Outliers make activation quantization difficult
- 3. Outliers persist in fixed channels.
- Offline migrates the quantization difficulty from activations to weights



(a) Original



(b) SmoothQuant

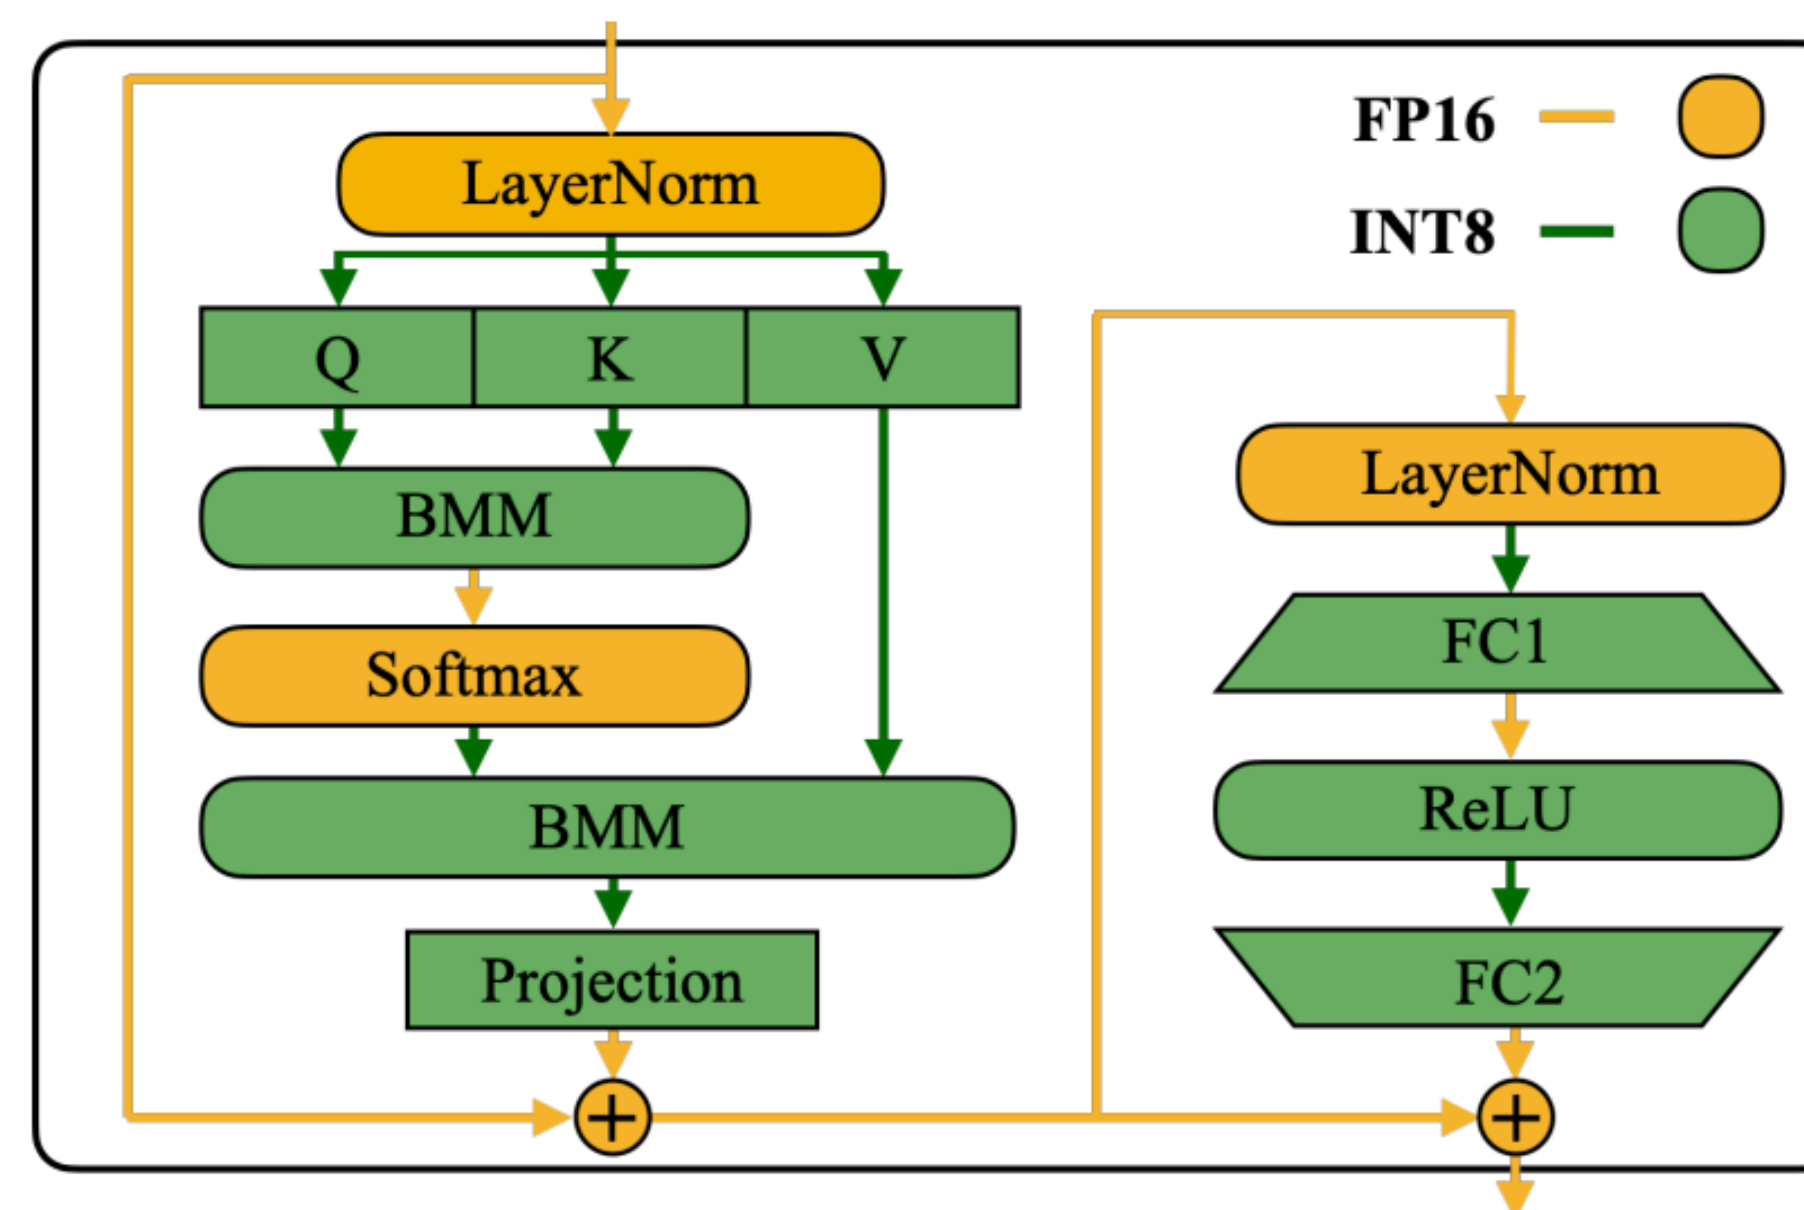

SmoothQuant (ICML 2023)

SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models

- “smooth” the input activation by dividing it by a per-channel smoothing factor

$$\mathbf{Y} = (\mathbf{X} \text{diag}(\mathbf{s})^{-1}) \cdot (\text{diag}(\mathbf{s}) \mathbf{W}) = \hat{\mathbf{X}} \hat{\mathbf{W}}$$


$$s_j = \max(|\mathbf{X}_j|)^\alpha / \max(|\mathbf{W}_j|)^{1-\alpha}$$



SmoothQuant (ICML 2023)

SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models

Method	OPT-175B	BLOOM-176B	GLM-130B*
FP16	71.6%	68.2%	73.8%
W8A8	32.3%	64.2%	26.9%
ZeroQuant	31.7%	67.4%	26.7%
LLM.int8()	71.4%	68.0%	73.8%
Outlier Suppression	31.7%	54.1%	63.5%
SmoothQuant-O1	71.2%	68.3%	73.7%
SmoothQuant-O2	71.1%	68.4%	72.5%
SmoothQuant-O3	71.1%	67.4%	72.8%



SmoothQuant (ICML 2023)

SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models





Figure 7: The PyTorch implementation of SmoothQuant-O3 achieves up to $1.51\times$ speedup and $1.96\times$ memory saving for OPT models on a single NVIDIA A100-80GB GPU, while `LLM.int8()` slows down the inference in most cases.

llama.cpp (15.2k stars on Github)

- 4-bit quantization / INT4-FP16 mixed precision
- Run on CPU
 - Run LLaMA-13B on 64GB M2 MacBook Pro
 - Run LLaMA-7B on 4GB RAM Raspberry Pi 4
- 10 sec/token

model	original size	quantized size (4-bit)
7B	13 GB	3.9 GB
13B	24 GB	7.8 GB
30B	60 GB	19.5 GB
65B	120 GB	38.5 GB



Artem Andreenko 
@miolini

...

I've sucefully runned LLaMA 7B model on my 4GB RAM Raspberry Pi 4. It's super slow about 10sec/token. But it looks we can run powerful cognitive pipelines on a cheap hardware.

```
937 → 'first'
767 → 'man'
373 → 'on'
278 → 'the'
18786 → 'moon'
471 → 'was'
29871 → ' '

sampling parameters: temp = 0.800000, top_k = 40, top_p = 0.950000, repeat_last_n = 64, repeat_penalty = 1.300000

The first man on the moon was 20 years old and looked like

top - 18:16:11 up 147 days, 9:22, 5 users, load average: 9.45, 8.06, 5.11
Tasks: 240 total, 2 running, 238 sleeping, 0 stopped, 0 zombie
%Cpu0 : 78.6 us, 7.8 sy, 0.0 ni, 0.0 id, 2.9 wa, 0.0 hi, 10.7 si, 0.0 st
%Cpu1 : 79.2 us, 13.2 sy, 0.0 ni, 0.0 id, 7.5 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu2 : 75.2 us, 16.2 sy, 0.0 ni, 0.0 id, 8.6 wa, 0.0 hi, 0.0 si, 0.0 st
%Cpu3 : 78.3 us, 13.2 sy, 0.0 ni, 0.0 id, 8.5 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 3792.3 total, 83.5 free, 3621.3 used, 87.5 buff/cache
MiB Swap: 65536.0 total, 60299.7 free, 5236.2 used, 46.4 avail Mem

  PID USER      PR  NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND
2705518 ubuntu    20   0 5231264 3.3g 1904 R 352.9  88.4  27:37.52 main
102 root      20   0      0     0     0 S 12.5   0.0  28:11.15 kswapd0

Features       : fp asimd evtstrm crc32 cpuid
CPU implementer : 0x41
CPU architecture: 8
CPU variant     : 0x0
CPU part        : 0xd08
CPU revision    : 3

Hardware       : BCM2835
Revision       : c03111
Serial         : 10000000d62b612e
Model          : Raspberry Pi 4 Model B Rev 1.1
ubuntu@rpi:~$
[0] 0:mc* 1:mc-
```

2:19 AM · Mar 13, 2023 · 1.4M Views

652 Retweets 196 Quotes 4,826 Likes

llama.cpp (15.2k stars on Github)



- Plain C/C++ implementation without dependencies
- 4-bit quantization for 99% normal values
- FP16 for 1% outliers
- Use CSC / CSR to mark the positions of outliers
- Very slow but simple and work method

OBQ (NIPS 2022)

Optimal Brain Compression: A Framework for Accurate Post-Training Quantization and Pruning

- Find a matrix of **quantized weights** \hat{W} which minimizes the squared error

$$\operatorname{argmin}_{\hat{\mathbf{W}}_\ell} \|\mathbf{W}_\ell \mathbf{X}_\ell - \hat{\mathbf{W}}_\ell \mathbf{X}_\ell\|_2^2$$

- Taylor approximation provides **explicit formulas for the optimal single weight to remove**, as well as the **optimal update of the remaining weights** which would compensate for the removal.

$$w_q = \operatorname{argmin}_{w_q} \frac{(\operatorname{quant}(w_q) - w_q)^2}{[\mathbf{H}_F^{-1}]_{qq}}, \quad \delta_F = -\frac{w_q - \operatorname{quant}(w_q)}{[\mathbf{H}_F^{-1}]_{qq}} \cdot (\mathbf{H}_F^{-1})_{:,q}.$$

OBQ (NIPS 2022)

Optimal Brain Compression: A Framework for Accurate Post-Training Quantization and Pruning

- OBQ handles each row independently in parallel, quantizing one weight at a time while always updating all not-yet-quantized weights, in order to compensate for the error incurred by quantizing a single weight

$$w_q = \operatorname{argmin}_{w_q} \frac{(\operatorname{quant}(w_q) - w_q)^2}{[\mathbf{H}_F^{-1}]_{qq}}, \quad \delta_F = -\frac{w_q - \operatorname{quant}(w_q)}{[\mathbf{H}_F^{-1}]_{qq}} \cdot (\mathbf{H}_F^{-1})_{:,q}.$$

- Gaussian elimination

$$\mathbf{H}_{-q}^{-1} = \left(\mathbf{H}^{-1} - \frac{1}{[\mathbf{H}^{-1}]_{qq}} \mathbf{H}_{:,q}^{-1} \mathbf{H}_{q,:}^{-1} \right)_{-p}.$$

OBB (NIPS 2022)

Optimal Brain Compression: A Framework for Accurate Post-Training Quantization and Pruning

Algorithm 1 Prune $k \leq d_{\text{col}}$ weights from row \mathbf{w} with inverse Hessian $\mathbf{H}^{-1} = (2\mathbf{X}\mathbf{X}^\top)^{-1}$ according to OBS in $O(k \cdot d_{\text{col}}^2)$ time.


$M = \{1, \dots, d_{\text{col}}\}$
for $i = 1, \dots, k$ **do**
 $p \leftarrow \operatorname{argmin}_{p \in M} \frac{1}{[\mathbf{H}^{-1}]_{pp}} \cdot w_p^2$
 $\mathbf{w} \leftarrow \mathbf{w} - \mathbf{H}_{:,p}^{-1} \frac{1}{[\mathbf{H}^{-1}]_{pp}} \cdot w_p$
 $\mathbf{H}^{-1} \leftarrow \mathbf{H}^{-1} - \frac{1}{[\mathbf{H}^{-1}]_{pp}} \mathbf{H}_{:,p}^{-1} \mathbf{H}_{p,:}^{-1}$
 $M \leftarrow M - \{p\}$
end for

- Total time complexity: $O(d_{\text{row}} \cdot d_{\text{col}}^3)$

GPTQ (ICLR 2023)

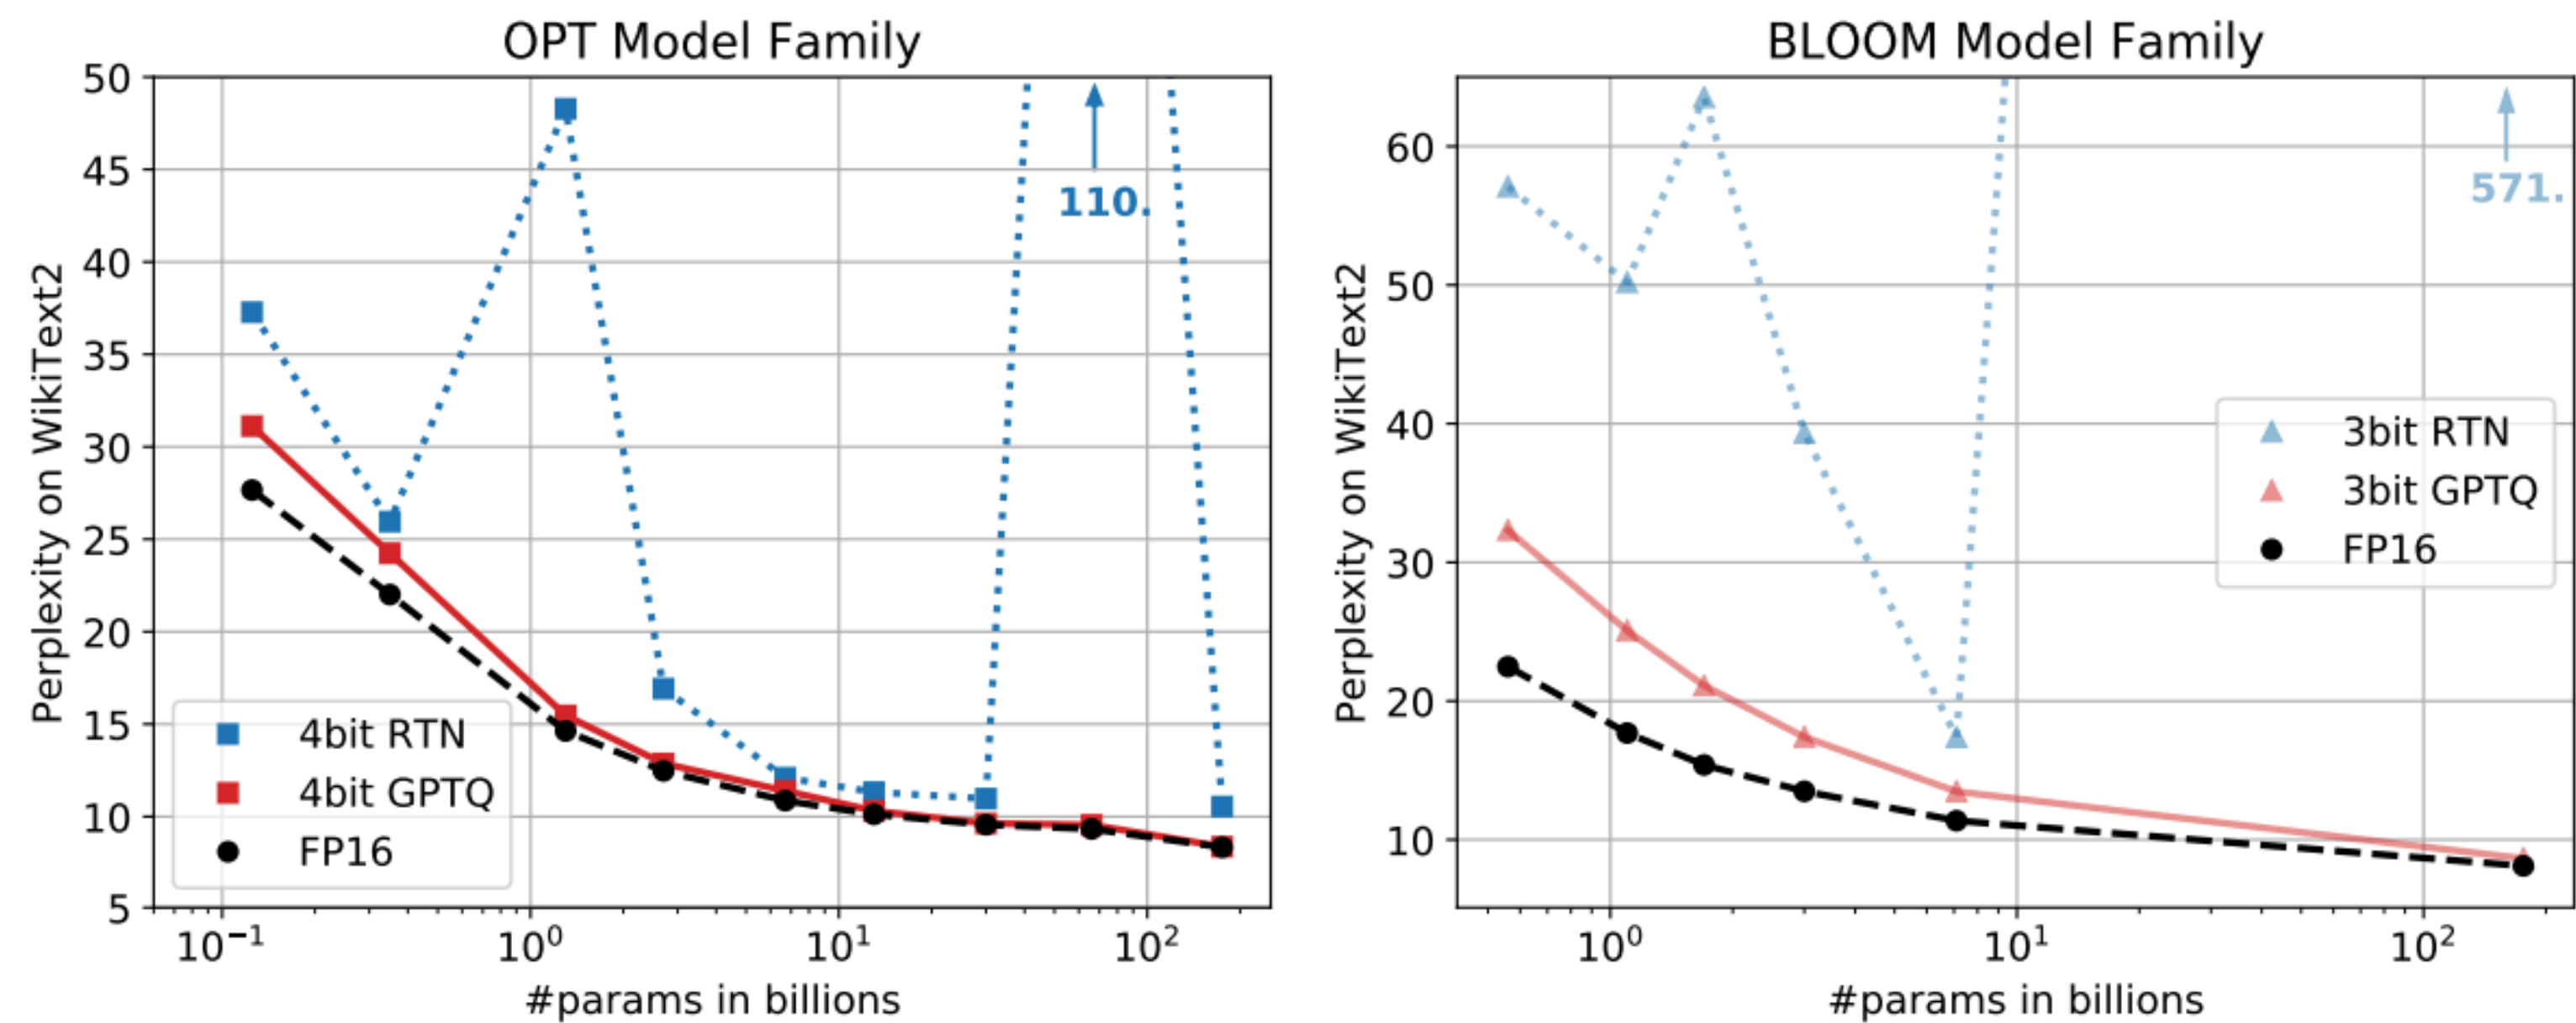
GPTQ: ACCURATE POST-TRAINING QUANTIZATION FOR GENERATIVE PRE-TRAINED

- Any fixed quantization order may perform well, especially on large models
- So we can quantize the weights of all rows in the same order
- Reduce the time complexity
from $O(d_{\text{row}} * d_{\text{col}}^3)$ to $O(\max\{d_{\text{row}} * d_{\text{col}}^2, d_{\text{row}}^3\})$



GPTQ (ICLR 2023)

GPTQ: ACCURATE POST-TRAINING QUANTIZATION FOR GENERATIVE PRE-TRAINED



OPT	Bits	125M	350M	1.3B	2.7B	6.7B	13B	30B	66B	175B
full	16	27.65	22.00	14.63	12.47	10.86	10.13	9.56	9.34	8.34
RTN	4	37.28	25.94	48.17	16.92	12.10	11.32	10.98	110	10.54
GPTQ	4	31.12	24.24	15.47	12.87	11.39	10.31	9.63	9.55	8.37
RTN	3	1.3e3	64.57	1.3e4	1.6e4	5.8e3	3.4e3	1.6e3	6.1e3	7.3e3
GPTQ	3	53.85	33.79	20.97	16.88	14.86	11.61	10.27	14.16	8.68

Table 3: OPT perplexity results on WikiText2.

Model	FP16	g128	g64	g32	3-bit
OPT-175B	8.34	9.58	9.18	8.94	8.68
BLOOM	8.11	9.55	9.17	8.83	8.64

Table 7: 2-bit GPTQ quantization results with varying group-sizes; perplexity on WikiText2.

GPTQ (ICLR 2023)

GPTQ: ACCURATE POST-TRAINING QUANTIZATION FOR GENERATIVE PRE-TRAINED

- GPT for LLaMA: <https://github.com/qwopqwop200/GPTQ-for-LLaMa>
- GPTQ does not improve the quantized performance on LLaMA (sometimes even worse)
- Trick: quantizing columns in order of decreasing activation size

LLaMA-65B	Bits	group-size	memory(MiB)	Wikitext2
FP16	16	-	OOM	3.53
RTN	4	-	-	3.92
GPTQ	4	-	OOM	3.84
GPTQ	4	128	OOM	3.65
RTN	3	-	-	10.59
GPTQ	3	-	OOM	5.04
GPTQ	3	128	OOM	4.17

Inference Acceleration

- Quantization
- Pruning
- Off-loading
- Distillation

SparseGPT

SparseGPT: Massive Language Models Can be Accurately Pruned in One-Shot

- Find a sparsity mask \mathbf{M} which minimizes the squared error

$$\operatorname{argmin}_{\text{mask } \mathbf{M}_\ell, \widehat{\mathbf{W}}_\ell} \|\mathbf{W}_\ell \mathbf{X}_\ell - (\mathbf{M}_\ell \odot \widehat{\mathbf{W}}_\ell) \mathbf{X}_\ell\|_2^2.$$

- The optimal values of all weights in the mask can be calculated exactly by solving the sparse reconstruction problem corresponding to each matrix row


$$w_p = \operatorname{argmin}_{w_p} \frac{w_p^2}{[\mathbf{H}^{-1}]_{pp}}, \quad \delta_p = -\frac{w_p}{[\mathbf{H}^{-1}]_{pp}} \cdot \mathbf{H}_{:,p}^{-1},$$

SparseGPT

SparseGPT: Massive Language Models Can be Accurately Pruned in One-Shot

Gaussian elimination:

$$(\mathbf{H}_{U_{j+1}})^{-1} = \left(\mathbf{B} - \frac{1}{[\mathbf{B}]_{11}} \cdot \mathbf{B}_{:,1} \mathbf{B}_{1,:} \right)_{2:,2:},$$



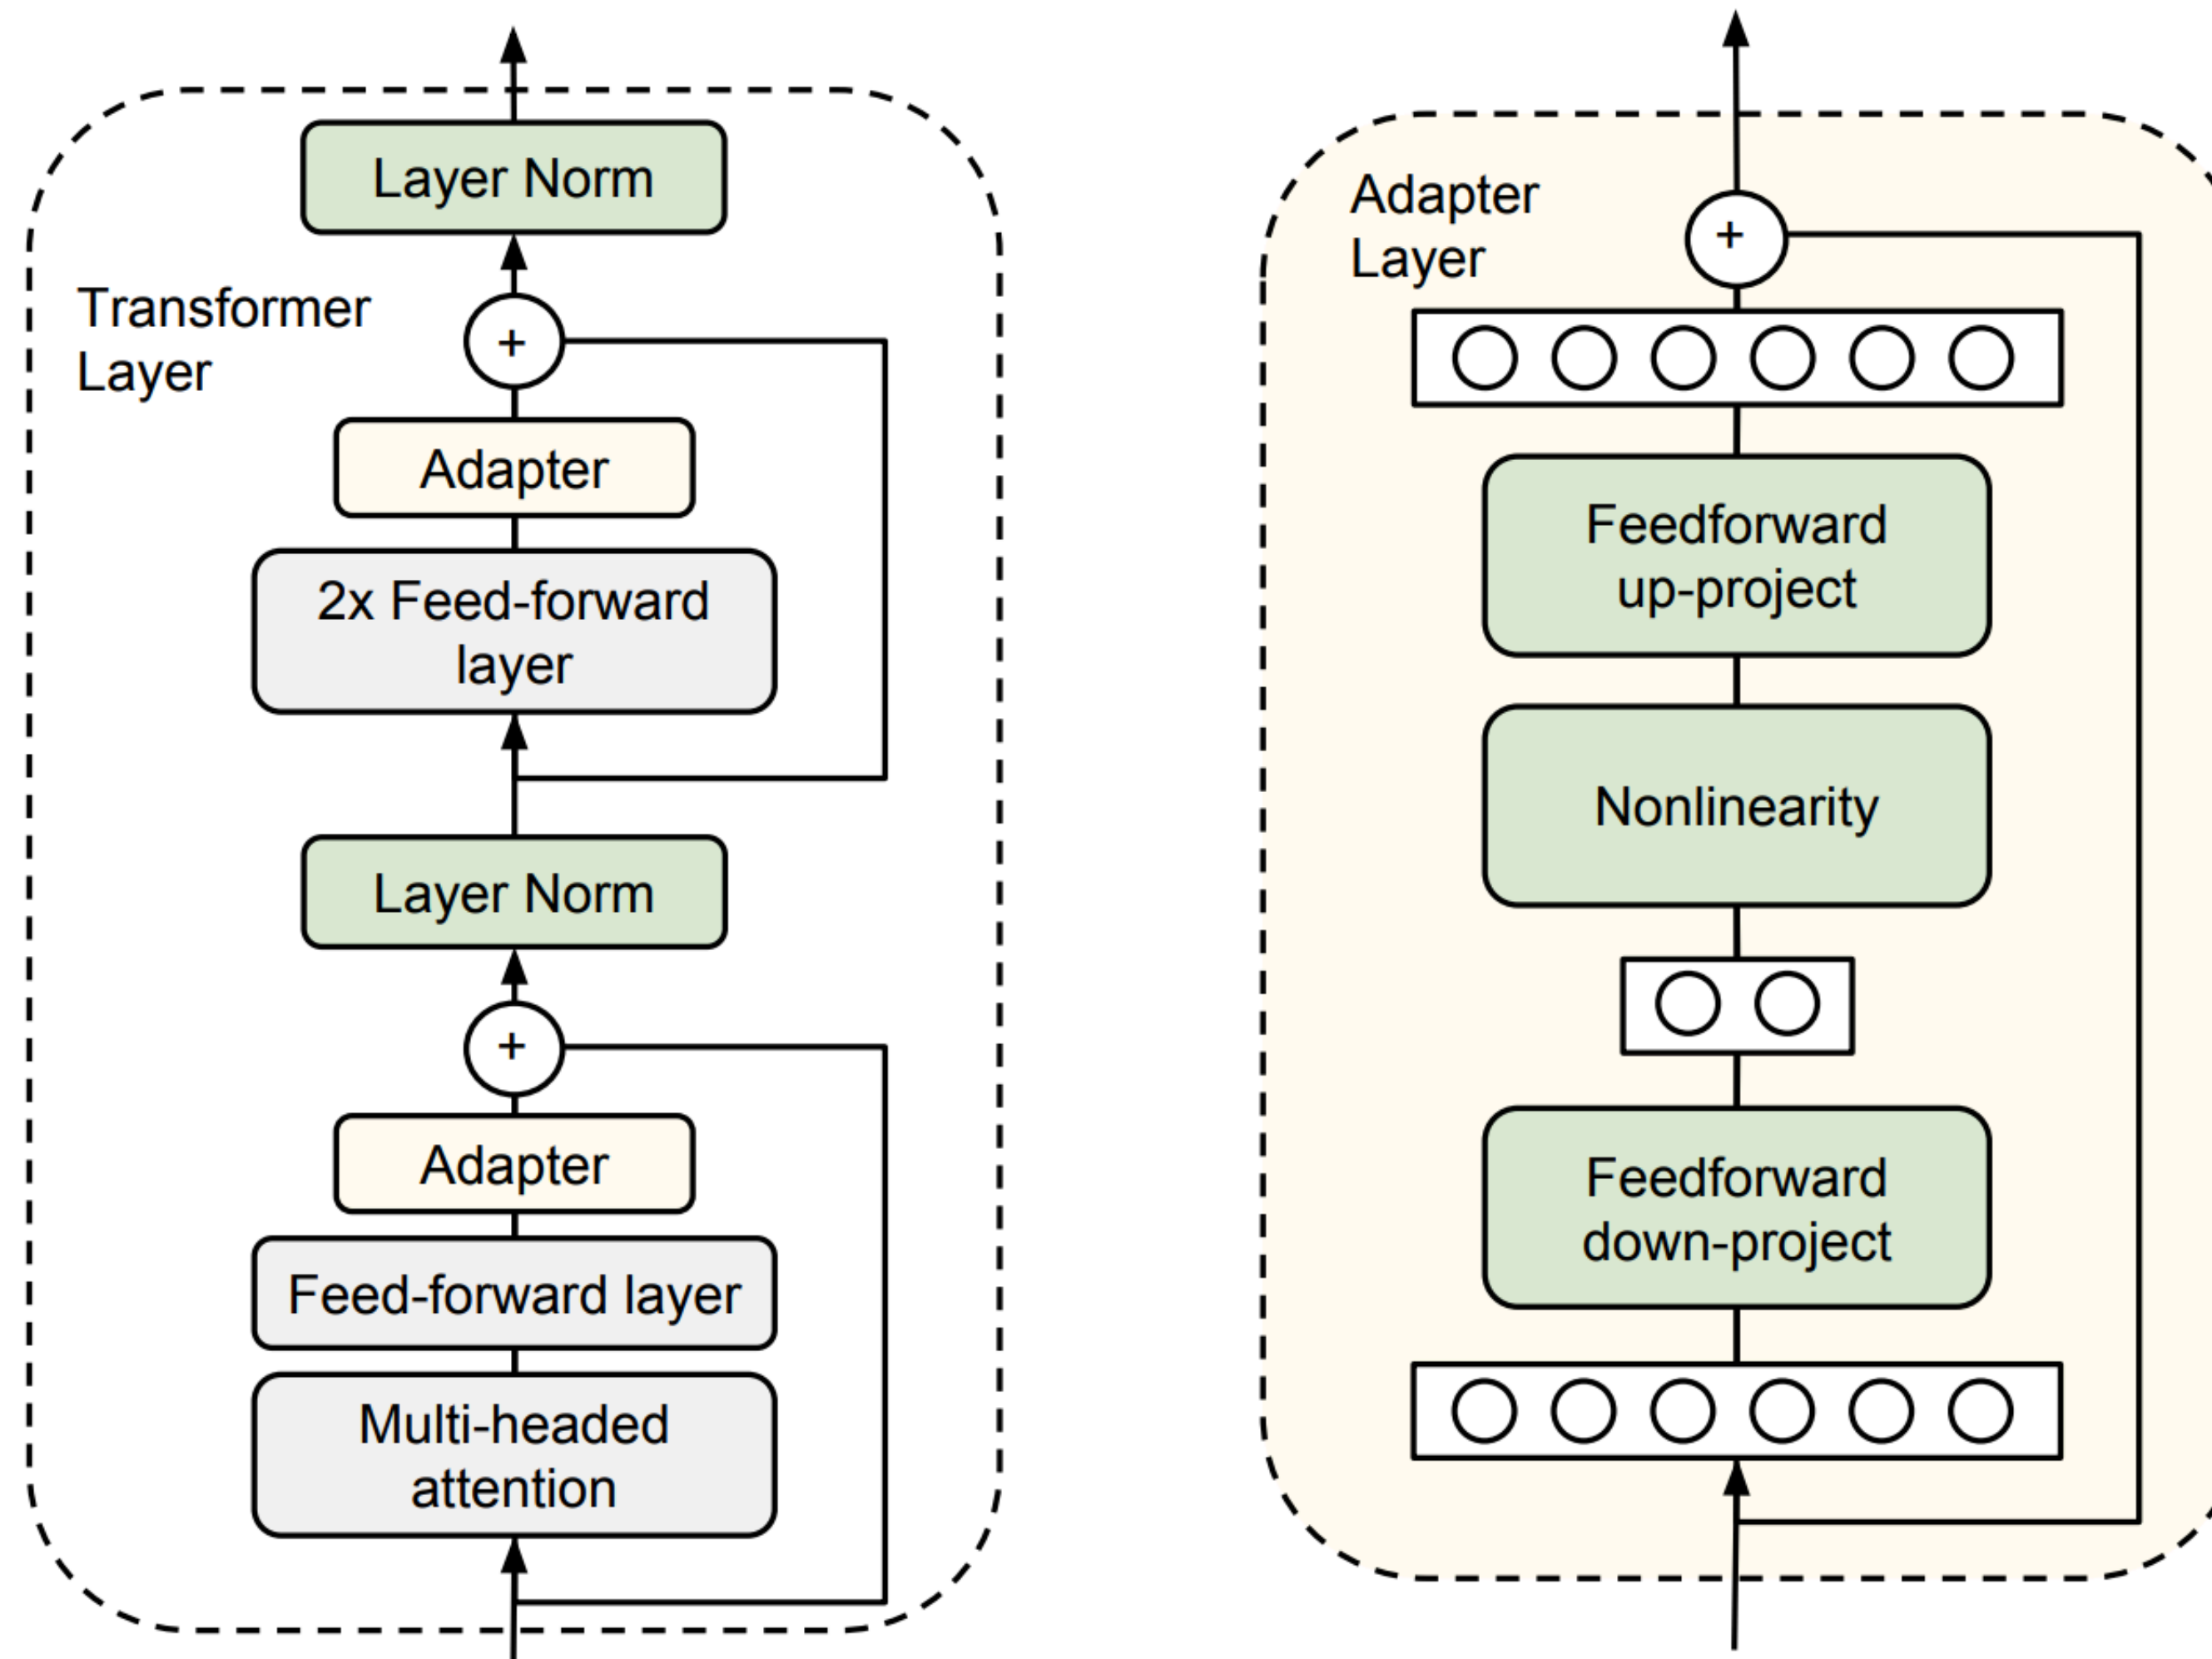
- Time Complexity: $O(d_{\text{row}} * d_{\text{col}}^2 + d_{\text{row}}^3)$

Fine-tuning

- Parameter Efficient Fine-tuning (PEFT)
 - Prompt Tuning (In-context Learning)
 - Prefix Tuning
 - P-Tuning V2 (ACL 2022)
- Adapter
 - Adapter tuning for NLP (ICML 2019)
 - Offsite-tuning (ICML 2023 submission)
- LoRA
 - Low Rank Adaptation for LLM



Adapter Tuning (ICML 2019)

P-Tuning v2: Prompt Tuning Can Be Comparable to Fine-tuning Universally Across Scales and Tasks



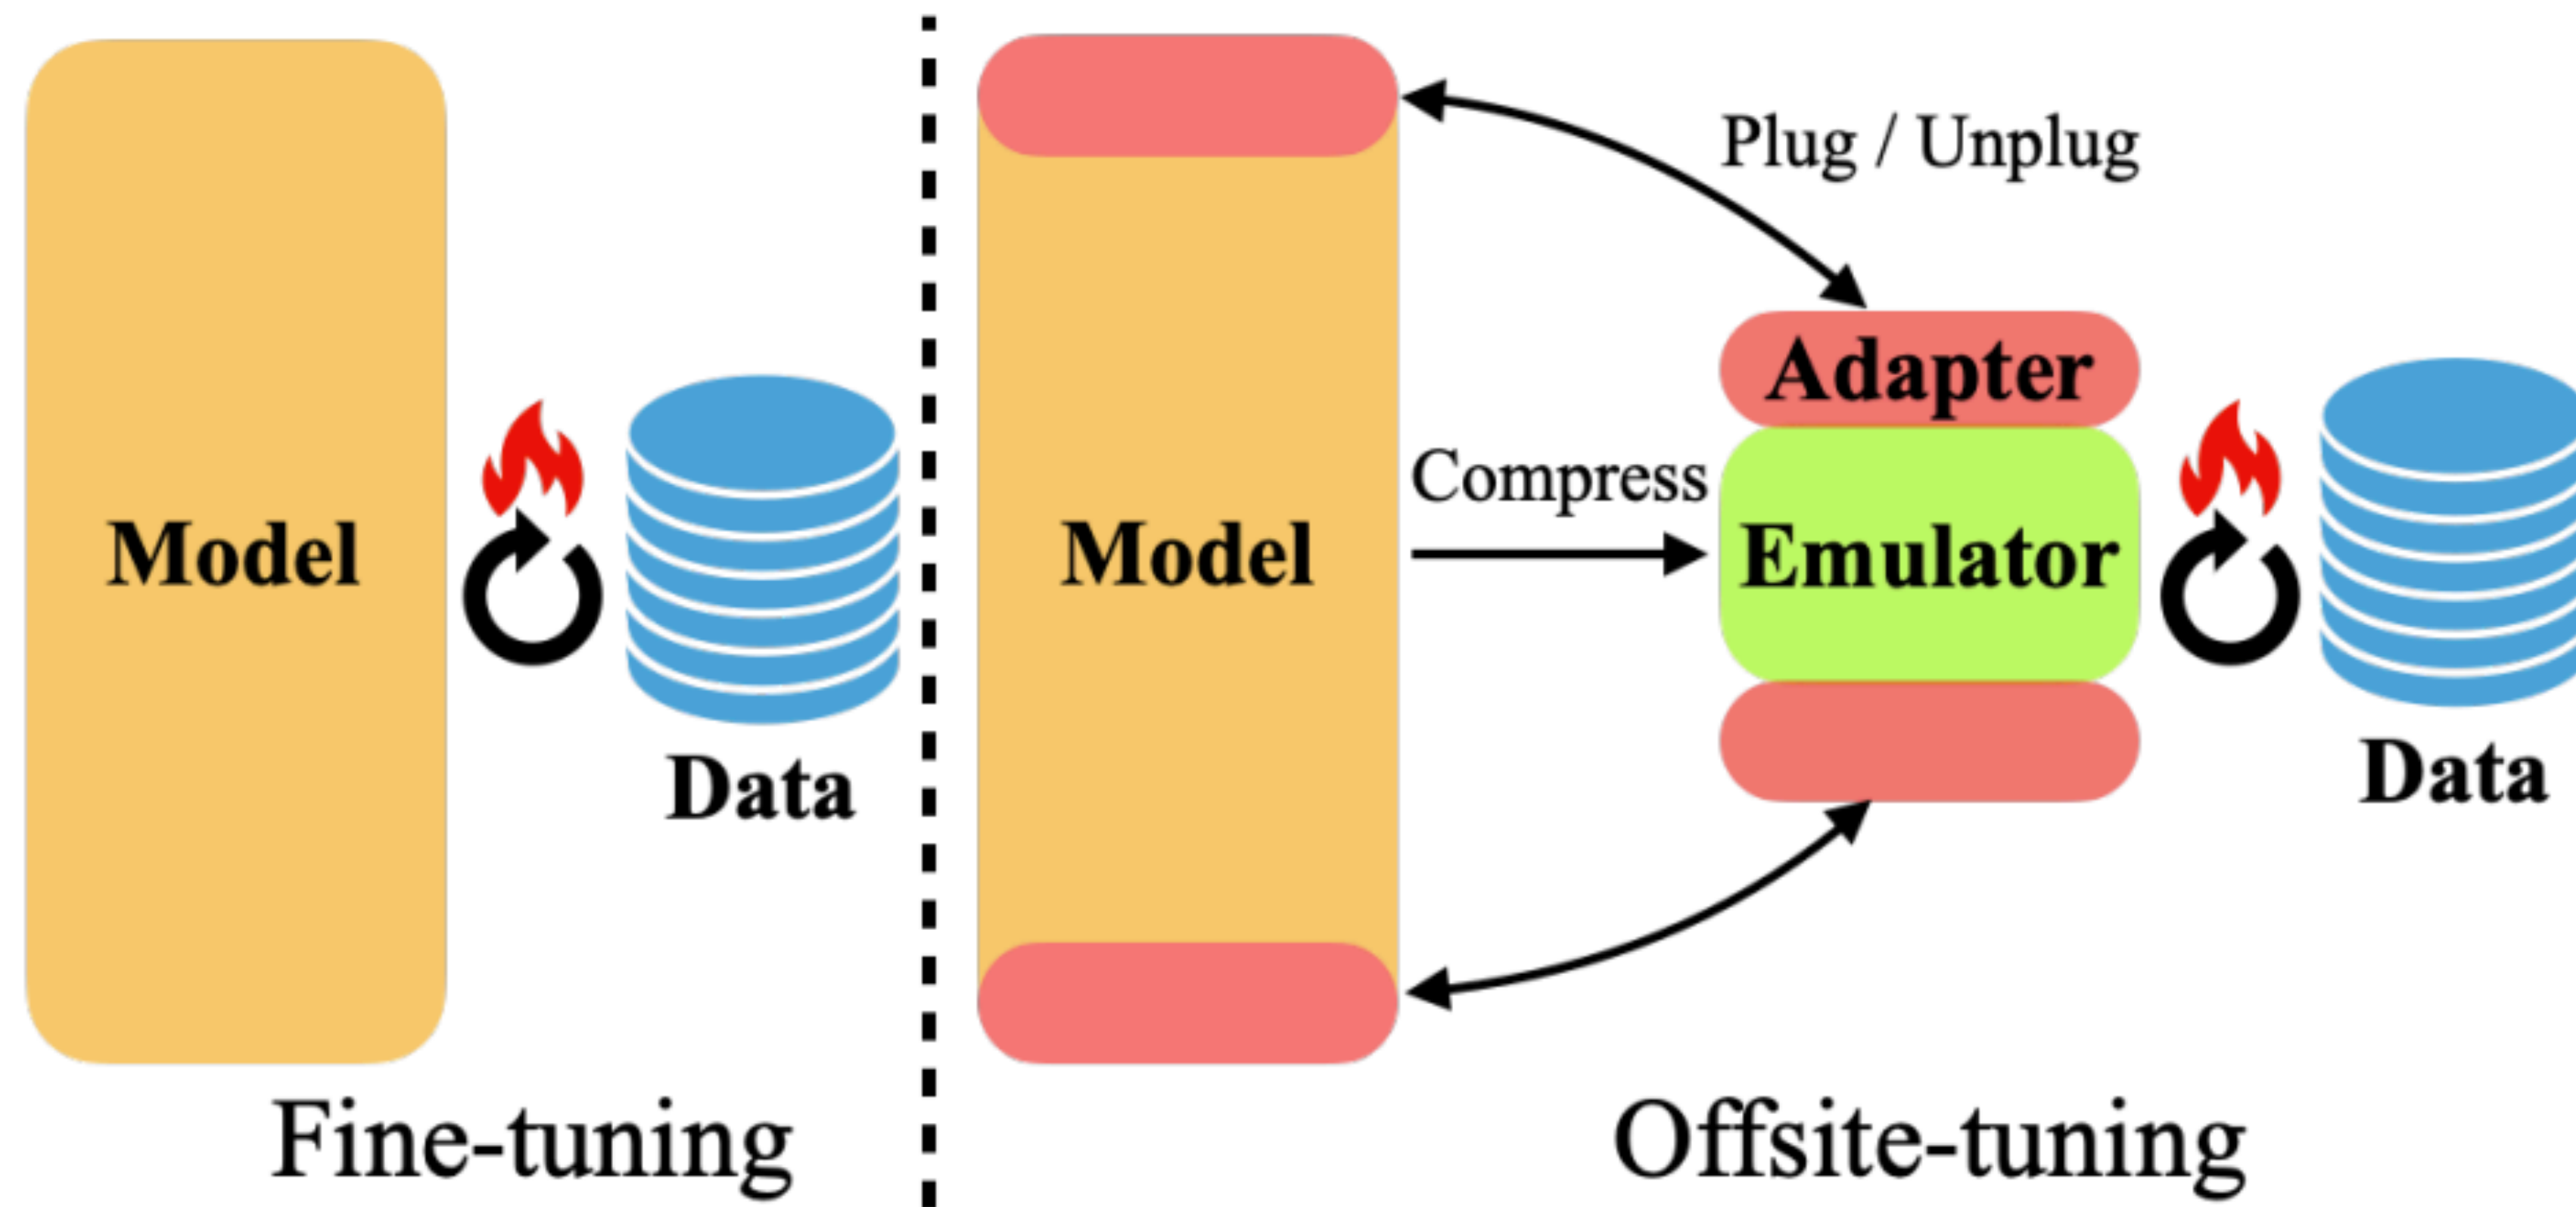
Adapter Tuning (ICML 2019)

P-Tuning v2: Prompt Tuning Can Be Comparable to Fine-tuning Universally Across Scales and Tasks



Offsite-Tuning (ICML 2023 submission)

Offsite-Tuning: Transfer Learning without Full Model



Offsite-Tuning (ICML 2023 submission)

Offsite-Tuning: Transfer Learning without Full Model

<i>Setting</i>	OpenBookQA	PIQA	ARC-E	ARC-C	HellaSwag	SciQ	WebQs	RACE	WikiText (↓)
GPT2-XL (2-16-2 Distill)									
Full ZS	23.0%	70.9%	58.2%	25.1%	40.0%	83.2%	1.5%	33.0%	20.44
Emulator ZS	18.8%	67.7%	53.2%	20.8%	33.5%	77.0%	0.2%	30.0%	25.12
FT	30.0%	73.2%	62.9%	30.0%	40.7%	92.5%	26.4%	43.2%	13.58
OT Emulator	24.0%	70.3%	58.2%	23.9%	35.8%	92.7%	18.9%	39.4%	17.64
OT Plug-in	28.2%	73.6%	61.4%	28.5%	41.6%	93.2%	19.9%	39.9%	14.94
OPT-1.3B (2-8-2 Distill)									
Full ZS	23.4%	71.6%	56.9%	23.5%	41.5%	84.4%	4.6%	34.2%	31.48
Emulator ZS	19.4%	68.7%	53.9%	21.5%	35.1%	80.9%	1.3%	33.0%	38.55
FT	31.4%	75.2%	61.3%	27.7%	42.7%	92.5%	31.2%	37.0%	12.52
OT Emulator	24.8%	71.6%	58.1%	26.1%	37.0%	92.2%	24.3%	38.6%	15.54
OT Plug-in	29.0%	74.5%	59.4%	27.8%	43.3%	92.9%	26.2%	38.9%	13.15

LoRA

LoRA: Low-Rank Adaptation of Large Language Models

$W_0 + \Delta W = W_0 + BA$, where $B \in \mathbb{R}^{d \times r}$, $A \in \mathbb{R}^{r \times k}$, and the rank $r \ll \min(d, k)$




Figure 1: Our reparametrization. We only train A and B .

LoRA

LoRA: Low-Rank Adaptation of Large Language Models

Model&Method	# Trainable Parameters	WikiSQL	MNLI-m	SAMSum
		Acc. (%)	Acc. (%)	R1/R2/RL
GPT-3 (FT)	175,255.8M	73.8	89.5	52.0/28.0/44.5
GPT-3 (BitFit)	14.2M	71.3	91.0	51.3/27.4/43.5
GPT-3 (PreEmbed)	3.2M	63.1	88.6	48.3/24.2/40.5
GPT-3 (PreLayer)	20.2M	70.1	89.5	50.8/27.3/43.5
GPT-3 (Adapter ^H)	7.1M	71.9	89.8	53.0/28.9/44.8
GPT-3 (Adapter ^H)	40.1M	73.2	91.5	53.2/29.0/45.1
GPT-3 (LoRA)	4.7M	73.4	91.7	53.8/29.8/45.9
GPT-3 (LoRA)	37.7M	74.0	91.6	53.4/29.2/45.1