

Modern DSLs Compile Workflow

Xinhao Luo

2025.11.28

Prerequisites

IR (Intermediate Representation)

- Def: Source Code → IR → Target Code
- Types:
 - High Level IR: Analysis and Transformation based on Source Code
 - Middle IR: Language- and Hardware-Independent Optimization
 - Low IR: Hardware-Specific Semantics and Optimizations
- Example: AST, ONNX Graph, Triton GPU IR

Prerequisites

Pass

- Def: Traversal over the IR to perform analysis or transformation.
- Types:
 - Analysis Pass: Read IR and collect information.
 - Transformation Pass: Modify IR to perform optimization or lowering.
- Example: Loop analysis, loop unroll

```
// 变换前: 未展开的循环
scf.for %i = %c0 to %c3 step %c1 {
    %val = arith.muli %i, %i : i32
    memref.store %val, %buffer[%i] : memref<4xi32>
}
```

loop unroll

```
// 变换后: 完全展开
%v0 = arith.muli %c0, %c0 : i32
memref.store %v0, %buffer[%c0] : memref<4xi32>
%v1 = arith.muli %c1, %c1 : i32
memref.store %v1, %buffer[%c1] : memref<4xi32>
%v2 = arith.muli %c2, %c2 : i32
memref.store %v2, %buffer[%c2] : memref<4xi32>
```

Prerequisites

Pattern Rewrite

- Def: Declarative mechanism to transform IR through local match + rewrite.
- Strategy: Patterns are applied **greedily** until the IR becomes stable.
- Example: Canonicalize, Operator Fusion

```
// 模式 1: x + 0 → x
%sum = arith.addi %x, %c0 : i32 // 变成 %x

// 模式 2: x * 1 → x
%prod = arith.muli %x, %c1 : i32 // 变成 %x

// 模式 3: x * 0 → 0
%prod = arith.muli %x, %c0 : i32 // 变成 %c0
```

```
// 初始 IR
%c0 = arith.constant 0 : i32
%c1 = arith.constant 1 : i32
%t1 = arith.addi %arg, %c0 : i32 // 可以简化
%t2 = arith.muli %t1, %c1 : i32 // 简化后又可以简化
func.return %t2 : i32
```

```
// 第一轮: 应用 x + 0 → x
%c1 = arith.constant 1 : i32
%t2 = arith.muli %arg, %c1 : i32 // 现在 %t1 变成了 %arg
func.return %t2 : i32
```

```
// 第二轮: 应用 x * 1 → x
func.return %arg : i32
```

Outline

- Modern DSLs Compile Workflow
 - TileLang
 - CuTeDSL, TLX, Triton
- Benchmark
- Conclusion

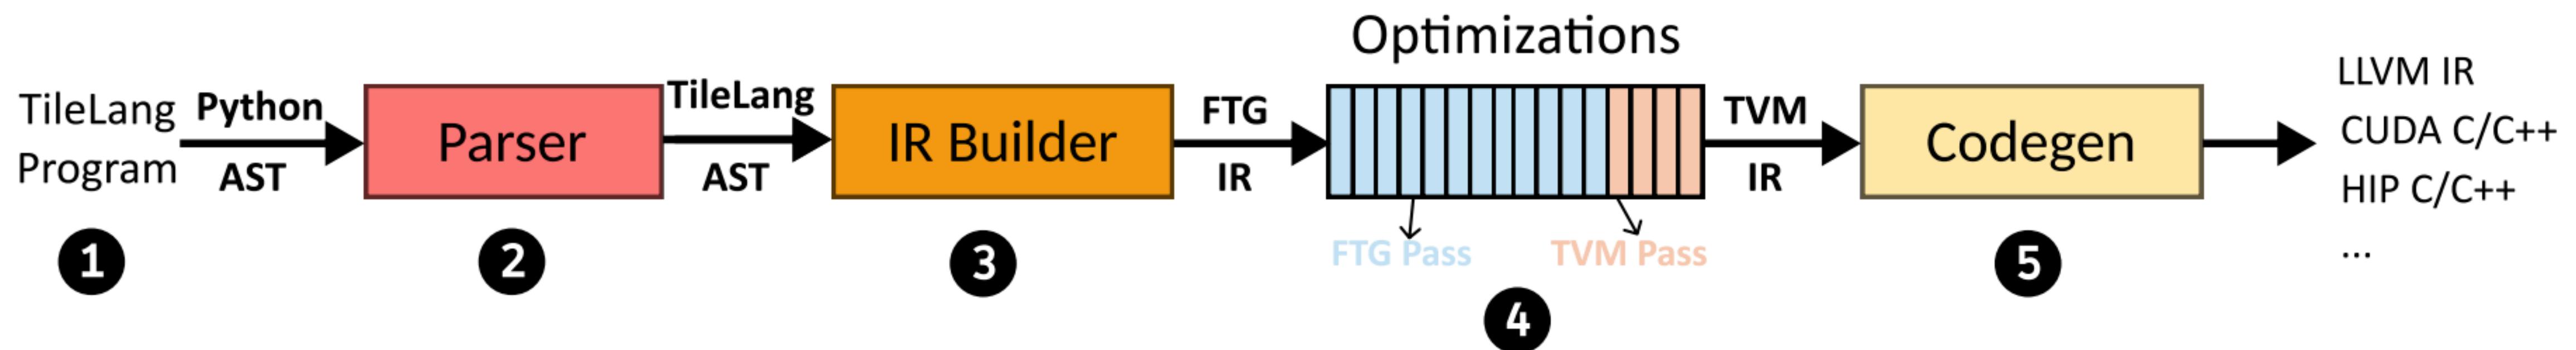


Modern DSLs

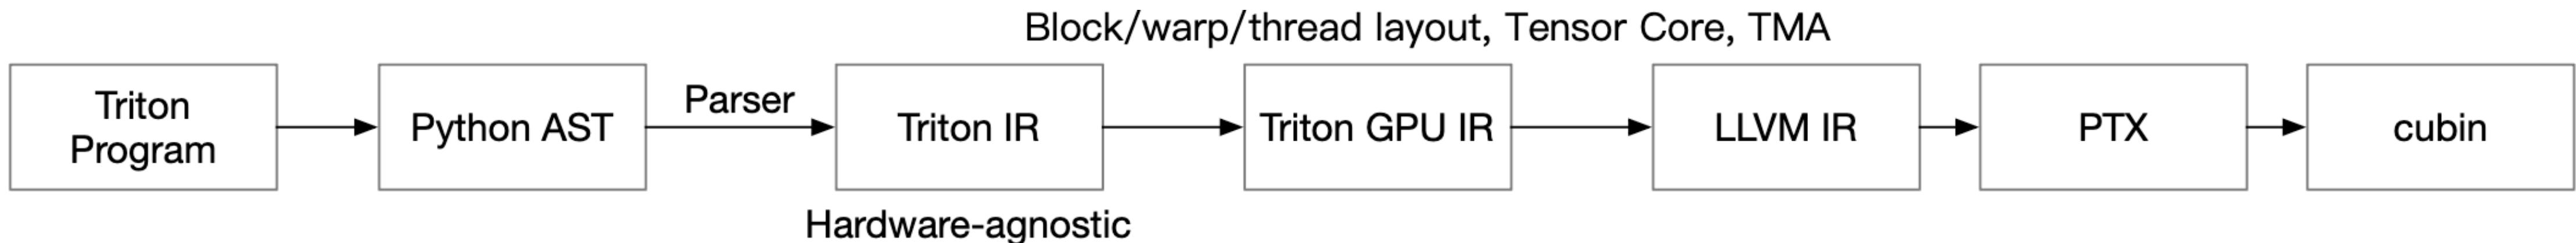
Overview

Workflow

- TileLang



- Triton



Key Points

Workflow

- Structure
- Pipeline
- **Pass: 作用以及Example**
- Optimization

Structure

Frontend

- Language
 - 显式 Grid/Block Level Tile Dataflow
 - 显式 Shared Memory/Local Register
 - TileOp: 高层算子抽象 (T.copy、T.gemm)
 - Software Pipeline

```
import tilelang.language as T

# 控制流与并行
with T.Kernel(T.ceildiv(M, 128), threads=256) as bx:
    with T.Parallel(128) as ty:
        with T.serial(16) as tx:
            ...

# 高层算子 (TileOp)
T.copy(src, dst)          # 数据搬移
T.gemm(A, B, C)          # 矩阵乘法
T.reduce_absmax(tensor)   # 归约操作

# 内存分配
buf = T.alloc_shared([128, 64], "float16")
acc = T.alloc_fragment([8, 8], "float32")

# 原子操作
T.atomic_add(ptr, val)
```

Structure Engine

- Phase
 - **两阶段 Pass Pipeline**
 - Phase-1: LowerAndLegalize, 把高层 DSL 变成可优化的 IR (TVM中的**TensorIR**)
 - Phase-2: OptimizeForTarget, 根据不同的目标硬件, 做不同的优化
- Lower:
 - Host/Device Codegen拆分与调度
 - Host IR → C++ wrapper
 - Device IR → CUDA/PTX

```
def matmul(M, N, K, block_M, block_N, block_K, dtype="float16",
           accum_dtype="float"):

    @T.prim_func
    def gemm(
        A: T.Tensor((M, K), dtype),
        B: T.Tensor((K, N), dtype),
        C: T.Tensor((M, N), dtype),
    ):
        with T.Kernel(T.ceildiv(N, block_N), T.ceildiv(M, block_M),
                      threads=128) as (bx, by):
            A_shared = T.alloc_shared((block_M, block_K), dtype)
            B_shared = T.alloc_shared((block_K, block_N), dtype)
            C_local = T.alloc_fragment((block_M, block_N), accum_)

            T.clear(C_local)
            for k in T.Pipelined(T.ceildiv(K, block_K), num_stages=2):
                T.copy(A[by * block_M, k * block_K], A_shared)
                T.copy(B[k * block_K, bx * block_N], B_shared)
                T.gemm(A_shared, B_shared, C_local)

            T.copy(C_local, C[by * block_M, bx * block_N])

    return gemm
```

Pipeline

Phase-1: LowerAndLegalize

- BindTarget
 - 将 Target (包含硬件能力信息) 绑定到 IR 中
 - 后续 Pass 需要查询硬件能力才能做决策

```
def matmul(M, N, K, block_M, block_N, block_K, dtype="float16",
           accum_dtype="float"):

    @T.prim_func
    def gemm(
        A: T.Tensor(M, K), dtype),
        B: T.Tensor(K, N), dtype),
        C: T.Tensor(M, N), dtype),
    ):
        with T.Kernel(T.ceildiv(N, block_N), T.ceildiv(M, block_M),
                      threads=128) as (bx, by):
            A_shared = T.alloc_shared((block_M, block_K), dtype)
            B_shared = T.alloc_shared((block_K, block_N), dtype)
            C_local = T.alloc_fragment((block_M, block_N), accum_dtype)

            T.clear(C_local)
            for k in T.Pipelined(T.ceildiv(K, block_K), num_stages=3):
                T.copy(A[by * block_M, k * block_K], A_shared)
                T.copy(B[k * block_K, bx * block_N], B_shared)
                T.gemm(A_shared, B_shared, C_local)

            T.copy(C_local, C[by * block_M, bx * block_N])

    return gemm
```

```
# 后续 Pass 会查询
if have_tma(target):    # Hopper 有 TMA 吗?
    use_tma_path()

if target.max_shared_memory_per_block >= 163840: # 共享内存够大吗?
    use_large_tile()
```

```
@I.ir_module
class Module:
    @T.prim_func
    def gemm(A_handle: T.handle, B_handle: T.handle, C_handle: T.handle):
        T.func_attr({"target": T.target({"arch": "sm_90", "host": {"keys": ["cpu"], "kind": "c", "tag": ""}, "keys": ["cuda", "gpu"], "kind": "cuda", "max_num_threads": 1024, "tag": "", "thread_warp_size": 32})})
```

Pipeline

Phase-1: LowerAndLegalize

- LetInline
 - 内联临时变量，减少嵌套深度
 - 让后续的 Pattern Rewrite 更容易
- Simplify
 - 常量折叠、代数化简、死代码消除

```
# 嵌套的Let表达式难以分析
Let x = a + b in
  (Let y = x * 2 in
    (y + x))

# LetInline处理后
(a + b) * 2 + (a + b)
```

```
# 简化前
if(True): a = 1
x = 0+ y * 1

# 简化后
a = 1
x = y
```

Pipeline

Phase-1: LowerAndLegalize

- InjectAssumes
 - 注入`_assume()`内置函数，告诉编译器“这个条件永远成立”
 - 消除冗余的边界检查，使用对齐的向量化加载 (`ldg.128`)

```
// 假设线程索引在合法范围内
__assume(threadIdx.x < blockDim.x);

// 假设指针满足特定对齐要求
__assume((global_ptr & 15) == 0); // 16字节对齐

// 假设循环次数满足特定条件
__assume(n % 8 == 0); // 向量化友好
```

```
# 没有对齐假设 → 保守方案
float4 value
value.x = ptr[0]
value.y = ptr[1]
value.z = ptr[2]
value.w = ptr[3]
# 生成：4条独立的32位加载指令

# 有对齐假设 → 激进优化
__assume((ptr & 15) == 0) # 保证16字节对齐
float4 value = *(float4*)ptr
# 生成：1条128位向量加载指令 ldg.128
```

Pipeline

Phase-1: LowerAndLegalize

- LayoutInference
 - 推断最终布局与线程映射，固化索引计算
 - 决定每个 Buffer 用什么布局，每个线程访问哪些元素，索引的线性化
 - 索引被改写为线性/规范形态

```
# 推断前 (抽象访问)
A_shared[ty, tx] = A_global[bx*128 + ty, tx]

# 推断后 (具体索引)
# 假设推断出 function 布局和特定线程映射
linear_idx = function(ty, tx, stride=128)
A_shared[linear_idx] = A_global[bx*128*K + ty*K + tx]
```

```
B_shared = T.Buffer((32, 128), "float16", scope="shared.dyn")
T.block_attr({"layout_map": {C_local: metadata["tl.Fragment"][0], A_shared: metadata["tl.Layout"][0], B_shared: metadata["tl.Layout"]
[1]}})

bx = T.launch_thread("blockIdx.x", 8)
by = T.launch_thread("blockIdx.y", 8)
tx = T.launch_thread("threadIdx.x", 128)
ty = T.launch_thread("threadIdx.y", 1)
tz = T.launch_thread("threadIdx.z", 1)
with T.block("tilelang_root"):
    T.reads(A[by * 128, 0:993], B[0:993, bx * 128], C[by * 128, bx * 128])
    T.writes()
    T.block_attr({"layout_map": {C_local: metadata["tl.Fragment"][0], A_shared: metadata["tl.Layout"][0], B_shared: metadata["tl.
Layout"]}[1]}})
    A_shared = T.alloc_buffer((128, 32), "float16", data=A_shared.data, scope="shared.dyn")
    B_shared = T.alloc_buffer((32, 128), "float16", data=B_shared.data, scope="shared.dyn")
    C_local = T.alloc_buffer((128, 128), data=C_local.data, scope="local.fragment")
    T.fill(T.tvm_access_ptr(T.type_annotation("float32"), C_local.data, 0, 16384, 2), 0)
```

Pipeline

Phase-1: LowerAndLegalize

- LowerTileOp
 - 把 TileOp (T.copy、T.gemm) 降级为底层 IR (循环、BufferLoad/Store)

```
# DSL
T.copy(A_global[0:128, 0:64], A_shared)

# 降级后 (伪代码)
for i in T.parallel(128):
    for j in T.vectorized(64):
        if (i < bound_i) and (j < bound_j): # 边界保护
            A_shared[i, j] = A_global[i, j]

# DSL
T.gemm(A, B, C)

# 降级后 (伪代码)
for i in range(M):
    for j in range(N):
        for k in range(K):
            C[i, j] += A[i, k] * B[k, j]
```

```
with T.block("tilelang_root"):
    T.reads(A[by * 128, 0:993], B[0:993, bx * 128], C[by * 128, bx * 128])
    T.writes()
    T.block_attr({"layout_map": {C_local: metadata["tl.Fragment"][0]}})

    A_shared = T.alloc_buffer((1, 16, 256), "float16", scope="shared.dyn")
    B_shared = T.alloc_buffer((2, 4, 512), "float16", scope="shared.dyn")
    C_local = T.alloc_buffer((128,), data=C_local.data, scope="local")

    if tx < 128 and tx >= 0:
        for i in T.unroll(64, annotations={"pragma_unroll_explicit": T.bool(False)}):
            for vec in T.vectorized(2):
                C_local[((i * 2 + vec) // 64 * 64 + tx // 32 * 16 + (i * 2 + vec) % 4 // 2 *
                         vec) % 64 // 4 * 8 + tx % 4 * 2 + (i * 2 + vec) % 2] = T.Cast("float32", 0)

        for k in T.serial(32, annotations={"num_stages": 3}):
            if tx == 0:
                T.tma_load(T.create_tma_descriptor(6, 2, A.data, 1024, 1024, T.int64(2), T.int64(
                    tvm_access_ptr(T.type_annotation("float16"), A_shared.data, 0, 4096, 2), k * 32,
                    if tx == 0:
                        for i in T.unroll(2):
                            T.tma_load(T.create_tma_descriptor(6, 2, B.data, 1024, 1024, T.int64(2), T.int64(
                                tvm_access_ptr(T.type_annotation("float16"), B_shared.data, 2048 * i, 2048,
                                T.tl_gemm("tl::gemm_ss<128, 128, 32, 4, 1, 0, 0, 0, 32, 128, 0, 0, true>", T.tvm_acce
                                A_shared.data, 0, 4096, 1), T.tvm_access_ptr(T.type_annotation("float16"), B_shared.d
                                type_annotation("float32"), C_local.data, 0, 16384, 3))
                                if tx < 128 and tx >= 0:
                                    for i in T.unroll(64, annotations={"pragma_unroll_explicit": T.bool(False)}):
                                        for vec in T.vectorized(2):
                                            C[by * 128 + (i * 2 + vec) // 64 * 64 + tx // 32 * 16 + (i * 2 + vec) % 4 //
                                                + vec) % 64 // 4 * 8 + tx % 4 * 2 + (i * 2 + vec) % 2] = T.Cast("float16", C_
                                                32 * 16 + (i * 2 + vec) % 4 // 2 * 8 + tx % 32 // 4) // 64 * 64 + ((i * 2 + v
                                                + vec) % 2) // 8 * 4 + ((i * 2 + vec) // 64 * 64 + tx // 32 * 16 + (i * 2 + v
                                                * 2 + ((i * 2 + vec) % 64 // 4 * 8 + tx % 4 * 2 + (i * 2 + vec) % 2) % 2])
```

Pipeline

Phase-1: LowerAndLegalize

- LegalizeVectorizedLoop
 - 修正向量化中的非法模式
- LegalizeSafeMemoryAccess
 - 为可能越界的访存注入保护

```
# 非法: 跨步向量化
for i in T.vectorized(8): # 步长=2
    A[i*2] = ...
# 修正后: 回退到标量
for i in T.serial(8):
    A[i*2] = ...
```

```
// 插入前
shared[idx] = val

// 插入后
if (idx < bound) {
    shared[idx] = val
} else {
    shared[idx] = 0
}
```

Pipeline

Phase-1: LowerAndLegalize

- Conclusion: 经过14个pass
 - 输入DSL → 规范TIR(循环、BufferLoad/Store, layout推断, 线性化索引, 向量化, 硬件信息绑定)

```

with T.block("root"):
    T.reads()
    T.writes()
    C_local = T.Buffer((128,), scope="local")
    T.block_attr({"layout_map": {C_local: metadata["tl.Fragment"][0]}})

    bx = T.launch_thread("blockIdx.x", 8)
    by = T.launch_thread("blockIdx.y", 8)
    tx = T.launch_thread("threadIdx.x", 128)
    ty = T.launch_thread("threadIdx.y", 1)
    tz = T.launch_thread("threadIdx.z", 1)

    with T.block("tilelang_root"):
        T.reads(A[by * 128, 0:993], B[0:993, bx * 128], C[by * 128, bx * 128])
        T.writes()
        T.block_attr({"layout_map": {C_local: metadata["tl.Fragment"][0]}})

        A_shared = T.alloc_buffer((1, 16, 256), "float16", scope="shared.dyn")
        B_shared = T.alloc_buffer((2, 4, 512), "float16", scope="shared.dyn")
        C_local = T.alloc_buffer((128,), data=C_local.data, scope="local")

        for i in T.unroll(64, annotations={"pragma_unroll_explicit": T.bool(False)}):
            for vec in T.vectorized(2):
                C_local[i * 2 + vec] = T.float32(0.0)

        for k in T.serial(32, annotations={"num_stages": 3}):
            if tx == 0:
                T.tma_load(T.create_tma_descriptor(6, 2, A.data, 1024, 1024, T.int64(2), T.int64(2048), 32, 128, 1, 1, 0, 2,
                                                tvm_access_ptr(T.type_annotation("float16"), A_shared.data, 0, 4096, 2), k * 32, by * 128, 0)
            if tx == 0:
                for i in T.unroll(2):
                    T.tma_load(T.create_tma_descriptor(6, 2, B.data, 1024, 1024, T.int64(2), T.int64(2048), 64, 32, 1, 1, 0,
                                                T.tvm_access_ptr(T.type_annotation("float16"), B_shared.data, i * 2048, 2048, 2), bx * 128 + i * 64, k)
            T.tl_gemm("tl::gemm_ss<128, 128, 32, 4, 1, 0, 0, 0, 32, 128, 0, 0, true>", T.tvm_access_ptr(T.type_annotation("float16"),
                A_shared.data, 0, 4096, 1), T.tvm_access_ptr(T.type_annotation("float16"), B_shared.data, 0, 4096, 1), T.tvm_access_ptr(T.type_annotation("float32"),
                C_local.data, 0, 16384, 3))

        for i in T.unroll(64, annotations={"pragma_unroll_explicit": T.bool(False)}):
            for vec in T.vectorized(2):
                C[by * 128 + i // 32 * 64 + tx // 32 * 16 + i % 2 * 8 + tx % 32 // 4, bx * 128 + i % 32 // 2 * 8 + tx % 4 *
                    Cast("float16", C_local[i * 2 + vec])]
```

Pipeline

Phase-2: OptimizeForTarget

- 根据目标硬件能力，走不同的特定优化路径
- TMA (Tensor Memory Accelerator) : bulk 异步复制
- Warp Specialization: Producer/Consumer 角色拆分
- WGMMA: 4-warp 组矩阵乘加
- mbarrier: 可计数字节屏障

Pipeline

Phase-2: OptimizeForTarget

- LowerSharedBarrier
 - 在共享内存中创建 mbarrier 并初始化

```
// 初始化
mbarrier_init(&barrier, num_threads);

// Producer 发起异步传输前
mbarrier_expect_tx(&barrier, bytes); // 告诉 barrier 期待 N 字节

// TMA 自动递减计数
tma_load_async(...);

// Producer 到达
mbarrier_arrive(&barrier);

// Consumer 等待
mbarrier_wait_parity(&barrier, parity); // parity ∈ {0, 1} 交替
```

```
A_shared = T.alloc_buffer((3, 1, 16, 256), "float16", scope="shared.dyn")
B_shared = T.alloc_buffer((3, 2, 4, 512), "float16", scope="shared.dyn")
C_local = T.alloc_buffer((128,), data=C_local.data, scope="local")
T.create_list_of_mbarrier(1, 1, 1, 128, 128, 128)
T.attr([128, 128], "kWarpSpecializationScope", 0)
if tx >= 128:
    for k in range(32):
        with T.block("", no_realize=True):
            T.reads()
            T.writes()
            T.block_attr({"stmt_group": 1})
            T.mbarrier_wait_parity(T.get_mbarrier(k % 3 + 3), T.bitwise_xor(k // 3 % 2, 1))
            if T.tl_shuffle_elect(128):
                T.tma_load(T.create_tma_descriptor(6, 2, A.data, 1024, 1024, T.int64(2), T.i
                0), T.get_mbarrier(k % 3), T.tvm_access_ptr(T.type_annotation("float16"), A_
                k * 32, by * 128, 0)
            with T.block("", no_realize=True):
                T.reads()
                T.writes()
                T.block_attr({"stmt_group": 1})
                if T.tl_shuffle_elect(128):
                    for i in T.unroll(2):
                        T.tma_load(T.create_tma_descriptor(6, 2, B.data, 1024, 1024, T.int64(2),
                        2, 0), T.get_mbarrier(k % 3), T.tvm_access_ptr(T.type_annotation("float16"),
                        * 4096, 2048, 2), bx * 128 + i * 64, k * 32, 0)
                T.ptx_arrive_barrier(T.get_mbarrier(k % 3))
```

Pipeline

Phase-2: OptimizeForTarget

- MultiVersionBuffer
 - 构造环形缓冲，支持生产/消费重叠

```
// 单版本 (无流水)
__shared__ float A[128][64]
for (int i = 0; i < N; i++) {
    copy(global, A);          // 阶段1: copy
    __syncthreads();
    compute(A);               // 阶段2: compute
    __syncthreads();
}

// 多版本 (3-阶段流水)
__shared__ float A[3][128][64];  // A[stage][...]
for (int i = 0; i < N; i++) {
    int stage = i % 3;
    copy(global, A[stage]);    // copy(i) 与 compute(i-1) 并发
    compute(A[(stage-1+3)%3]);
}
```

```
T.block_attr({"layout_map": {C_local: metadata["tl.Fragment"][0]}})
A_shared = T.alloc_buffer((1, 16, 256), "float16", scope="shared.dyn")
B_shared = T.alloc_buffer((2, 4, 512), "float16", scope="shared.dyn")
C_local = T.alloc_buffer((128,), data=C_local.data, scope="local")

A_shared = T.alloc_buffer((3, 1, 16, 256), "float16", scope="shared.dyn")
B_shared = T.alloc_buffer((3, 2, 4, 512), "float16", scope="shared.dyn")
C_local = T.alloc_buffer((128,), data=C_local.data, scope="local")
```

Pipeline

Phase-2: OptimizeForTarget

- WarpSpecialized
 - Producer 在加载 tile(i+1) 时, Consumer 在计算 tile(i)
 - 寄存器压力降低: Producer不需要accumulator, Consumer 不需要地址计算

```
__global__ void kernel() {
    for (int i = 0; i < N; i++) {
        // 所有 warp 都执行相同的序列
        copy_phase();      // copy
        __syncthreads();
        compute_phase();   // compute
        __syncthreads();
    }
}

__global__ void kernel() {
    int warp_id = threadIdx.x / 32;

    if (warp_id < 4) {
        // Producer warpgroup (warp 0-3)
        for (int i = 0; i < N; i++) {
            mbarrier_expect_tx(&barrier[i%3], bytes);
            tma_load_async(global, shared[i%3]); // 只负责 copy
            mbarrier_arrive(&barrier[i%3]);
        }
    } else {
        // Consumer warpgroup (warp 4-7)
        for (int i = 0; i < N; i++) {
            mbarrier_wait_parity(&barrier[i%3], i%2);
            wmma(shared[i%3], acc); // 只负责 compute
        }
    }
}
```

```
if tx >= 128:
    for k in range(32):
        with T.block("", no_realize=True):
            T.reads()
            T.writes()
            T.block_attr({"stmt_group": 1})
            T.mbarrier_wait_parity(T.get_mbarrier(k % 3 + 3), T.bitwise_xor(k // 3 % 2, 1))
            if T.tl_shuffle_elect(128):
                T.tma_load(T.create_tma_descriptor(6, 2, A.data, 1024, 1024, T.int64(2), T.i
                    0), T.get_mbarrier(k % 3), T.tvm_access_ptr(T.type_annotation("float16"), A_
                    k * 32, by * 128, 0))
            with T.block("", no_realize=True):
                T.reads()
                T.writes()
                T.block_attr({"stmt_group": 1})
                if T.tl_shuffle_elect(128):
                    for i in T.unroll(2):
                        T.tma_load(T.create_tma_descriptor(6, 2, B.data, 1024, 1024, T.int64(2),
                            2, 0), T.get_mbarrier(k % 3), T.tvm_access_ptr(T.type_annotation("float1
                            * 4096, 2048, 2), bx * 128 + i * 64, k * 32, 0))
                    T.ptx_arrive_barrier(T.get_mbarrier(k % 3))
                else:
                    for i in T.unroll(64, annotations={"pragma_unroll_explicit": T.bool(False)}):
                        for vec in T.vectorized(2):
                            C_local[i * 2 + vec] = T.float32(0.0)
                    for k in range(32):
                        with T.block("", no_realize=True):
                            T.reads()
                            T.writes()
                            T.block_attr {"stmt_group": 1}
                            T.mbarrier_wait_parity(T.get_mbarrier(k % 3), k // 3 % 2)
                            T.tl_gemm("tl::gemm_ss<128, 128, 32, 4, 1, 0, 0, 0, 32, 128, 0, 0, true>", T.tvm
                                "float16"), A_shared.data, k % 3 * 4096, 4096, 1), T.tvm_access_ptr(T.type_anno
```

Pipeline

Phase-2: OptimizeForTarget

- 输入 Phase-1 的规范 IR:

```
// 伪 IR
for tile in range(N):
    for thread in parallel:
        if in_copy_region:
            shared[idx] = global[idx] // copy
            __syncthreads()

        for thread in parallel:
            if in_compute_region:
                acc += shared[idx] * weight[idx] // compute
                __syncthreads()
```

- 输出:

```
// 伪 IR
int warp_role = get_warp_role(); // Producer/Consumer/Both

if (warp_role == Producer) {
    for tile in range(N):
        mbarrier_expect_tx(&barrier[tile%stages], bytes);
        tma_load_async(global, shared[tile%stages]);
        mbarrier_arrive(&barrier[tile%stages]);
}

if (warp_role == Consumer) {
    for tile in range(N):
        mbarrier_wait_parity(&barrier[tile%stages], tile%2);
        wmma(shared[tile%stages], acc);
}

if (warp_role == Both) {
    // epilogue
    ...
}
```

- Pass 的改写步骤:

- 1. 找到 copy 和 compute 的分界
- 2. 为循环/分支打标签:
 - Producer: copy
 - Consumer: compute
 - Both: epilogue
- 3. 插入 mbarrier
 - Producer: expect_tx → tma → arrive
 - Consumer: wait → compute
- 4. 分裂控制流: 用 if (warp_role) 条件化执行

Pipeline

Phase-2: OptimizeForTarget

- AnnotateWarpGroupRegAlloc
 - 为 WS 的不同分支注入寄存器预算设置
- PipelinePlanning
 - 计算流水线的阶段数、调度顺序、依赖关系

```
if 128 <= tx:  
    T.set_max_nreg(24, 0)  
    for k in range(32):  
        T.mbarrier_wait_parity(T.ge  
    if T.tl_shuffle_elect(128):  
        T.mbarrier_expect_tx(T.  
        T.fence_proxy_async()  
        T.tma_load(A_desc, T.ge  
        4096, 2), k * 32, by *  
        T.mbarrier_expect_tx(T.  
        T.fence_proxy_async()  
        T.tma_load(B_desc, T.ge  
        2), bx * 128, k * 32, 0  
        T.tma_load(B_desc, T.ge  
        2048, 2), bx * 128 + 64  
        T.ptx_arrive_barrier(T.get_  
  
    else:  
        T.set_max_nreg(240, 1)
```

```
# 注解示例 (在 IR 上)  
loop.attr["software_pipeline_stage"] = [0, 1, 2]      # 阶段 ID  
loop.attr["software_pipeline_order"] = [0, 2, 1]      # 执行顺序  
loop.attr["software_pipeline_async_stages"] = 2       # 异步阶段数
```

- InjectSoftwarePipeline
 - 根据规划，展开为实际的 prologue/mainloop/epilogue。

```
// 原始循环  
for (int i = 0; i < N; i++) {  
    stage0: mbarrier_expect_tx + tma_load  
    stage1: mbarrier_wait + compute  
}  
  
// 展开后 (3-阶段)  
// === Prologue ===  
stage0(tile=0); // 预加载  
stage0(tile=1);  
  
// === Main Loop ===  
for (int i = 0; i < N-2; i++) {  
    stage1(tile=i); // compute tile i  
    stage0(tile=i+2); // load tile i+2  
}  
  
// === Epilogue ===  
stage1(tile=N-2);  
stage1(tile=N-1);
```

Pipeline

Phase-2: OptimizeForTarget

- RewriteWgmmaSync
 - Warp-Group Matrix-Multiply-Accumulate, Hopper 的 Tensor Core 指令。
 - 特点：
 - 1. 以 4 个 warp 为单位执行异步，发射后立即返回
 - 2. 需要显式同步：`warpgroup_wait<N>`
 - `warpgroup_wait<N>` 语义：
 - N=0：等待所有未完成的 WGMMA； N=1：等待除最近 1 条外的所有 WGMMA
 - 收集同一作用域内的所有 WGMMA，在最后一个之前插入 `warpgroup_wait<N>`，合并多余的等待语句

```
// 原始 (无同步)
wgmma.mma_async(...); // 发射 WGMMA 1
wgmma.mma_async(...); // 发射 WGMMA 2
// ... 立即使用 accumulator

// 改写 (插入组同步)
wgmma.mma_async(...); // 发射 WGMMA 1
wgmma.mma_async(...); // 发射 WGMMA 2
warpgroup_wait<0>; // 等待所有 WGMMA 完成
// ... 安全使用 accumulator
```

Pipeline

Phase-2: OptimizeForTarget

- Conclusion: 经过13个pass
 - 创建了 mbarrier
 - 构造了缓冲 (环形阶段)
 - 拆分了 Producer/Consumer (Warp Specialization)
 - 注入了 TMA (expect_tx/arrive/wait_parity)
 - 规划并展开了流水线 (prologue/mainloop/epilogue)
 - 改写了 WGMMA 同步 (warpgroup_wait)

```
@T.prim_func
def gemm_kernel(A_desc: T.handle("uint8x128", "grid_constant"), B_desc: T.handle("uint8x128", "grid_constant"), C: T.handle("float16", "global")):
    T.func_attr({"calling_conv": 2, "dyn_shared_memory_buf": 49152, "target": T.target({"arch": "sm_90", "keys": ["cuda", "gpu"], "kind": "cuda", "max_num_threads": 1, "threadIdx.x.z": 8, "blockIdx.y": 8, "threadIdx.x": 256, "threadIdx.y": 1, "tir.noalias": True})}
    C_1 = T.decl_buffer((1048576,), "float16", data=C)
    C_local = T.handle("float32", "local")
    C_local_1 = T.decl_buffer((128,), data=C_local, scope="local")
    bx = T.launch_thread("blockIdx.x", 8)
    buf_dyn_shmem = T.allocate([49152], "uint8", "shared.dyn")
    C_local = T.allocate([128], "float32", "local")
    by = T.launch_thread("blockIdx.y", 8)
    tx = T.launch_thread("threadIdx.x", 256)
    T.create_barriers(6)
    if T.tl_shuffle_select(0):
        T.call_extern("handle", "tl::prefetch_tma_descriptor", A_desc)
        T.call_extern("handle", "tl::prefetch_tma_descriptor", B_desc)
        T.ptx_init_barrier_thread_count(T.get_mbarrier(0), 128)
        T.ptx_init_barrier_thread_count(T.get_mbarrier(1), 128)
        T.ptx_init_barrier_thread_count(T.get_mbarrier(2), 128)
        T.ptx_init_barrier_thread_count(T.get_mbarrier(3), 128)
        T.ptx_init_barrier_thread_count(T.get_mbarrier(4), 128)
        T.ptx_init_barrier_thread_count(T.get_mbarrier(5), 128)
        T.ptx_fence_barrier_init()
        T.tvm_storage_sync("shared")
        ty = T.launch_thread("threadIdx.y", 1)
        tz = T.launch_thread("threadIdx.z", 1)
        T.tarr([128, 128], "WarpSpecializationScope", 0)
        if 128 <= tx:
            T.set_max_nreg(24, 0)
            for k in range(32):
                T.mbarrier_wait_parity(T.get_mbarrier(k % 3 + 3), T.bitwise_xor(k % 6 // 3, 1))
                if T.tl_shuffle_select(128):
                    T.mbarrier_expect_tx(T.get_mbarrier(k % 3), 8192)
                    T.fence_proxy_async()
                    T.tma_load(A_desc, T.get_mbarrier(k % 3), T.tvm_access_ptr(T.type_annotation("float16"), buf_dyn_shmem, 12288 + k % 3 * 4096, 4096, 2), k * 32, by * 1)
                    T.mbarrier_expect_tx(T.get_mbarrier(k % 3), 8192)
                    T.fence_proxy_async()
                    T.tma_load(B_desc, T.get_mbarrier(k % 3), T.tvm_access_ptr(T.type_annotation("float16"), buf_dyn_shmem, k % 3 * 4096, 2048, 2), bx * 128, k * 32, 0)
                    T.tma_load(B_desc, T.get_mbarrier(k % 3), T.tvm_access_ptr(T.type_annotation("float16"), buf_dyn_shmem, k % 3 * 4096 + 2048, 2048, 2), bx * 128 + 64, 0)
                    T.ptx_arrive_barrier(T.get_mbarrier(k % 3))
            else:
                T.set_max_nreg(240, 1)
                for i in T.unroll(64):
                    C_local_1[i * 2:i * 2 + 2] = T.Broadcast(T.float32(0.0), 2)
                for k in range(32):
                    T.mbarrier_wait_parity(T.get_mbarrier(k % 3), k % 6 // 3)
                    T.fence_proxy_async()
                    T.tl_gemm("tl::gemm_ss<128, 128, 32, 4, 1, 0, 0, 32, 128, 0, 0, true>", T.tvm_access_ptr(T.type_annotation("float16"), buf_dyn_shmem, 12288 + k % 3 * 4096, 4096, 1), T.tvm_access_ptr(T.type_annotation("float32"), C.local, 0, 16384, 3))
                    T.ptx_arrive_barrier(T.get_mbarrier(k % 3 + 3))
                for i in T.unroll(64):
                    C_1[i * 131072 + i // 32 * 65536 + tx // 32 * 16384 + i % 2 * 8192 + tx % 32 // 4 * 1024 + bx * 128 + i % 32 // 2 * 8 + tx % 4 * 2:by * 131072 + i // 32 * 1024 + bx * 128 + i % 32 // 2 * 8 + tx % 4 * 2 + 2] = T.Cast("float16x2", C_local_1[i * 2:i * 2 + 2])
    
```

Benchmarks

Benchmark Results

Dense GEMM & Attention

- **H100 SXM5 80GB, FP16**
- **GEMM: MNK = 8192**
- **Attention: Seqlen = 4096, Heads = 32, Dim = 128**
- 参考性能，来源是开源仓库中的example

	GEMM	Attention
CUTLASS	830	602
CuTeDSL	806	650
TileLang	743	606
TLX	665	594
Triton	635	515

Conclusion

Conclusion

Just use DSLs to write kernels

维度	CuTeDSL	TileLang	Triton	TLX
易用性	较低，需要深入理解CUTLASS概念	中等	较高	中等 基于Triton改进
性能上限	最高 匹配手写CUDA性能	较高	中等 通用性好但峰值受限	中等偏上 针对性优化
优化表达空间	最丰富 可控制底层硬件细节	丰富 Pass配置+硬件特定优化	适中 高级抽象易用但限制多	较丰富 扩展Triton优化能力
硬件适配性	专门为NVIDIA优化	良好 多后端支持	良好 支持多硬件平台	专注NVIDIA特性的扩展
生态成熟度	NVIDIA官方维护	中等 快速发展中	最高 社区活跃用户多	较低

Thanks