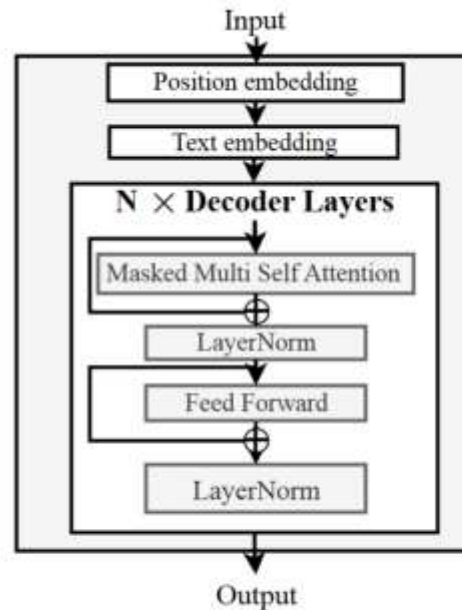


# Speculative Decoding

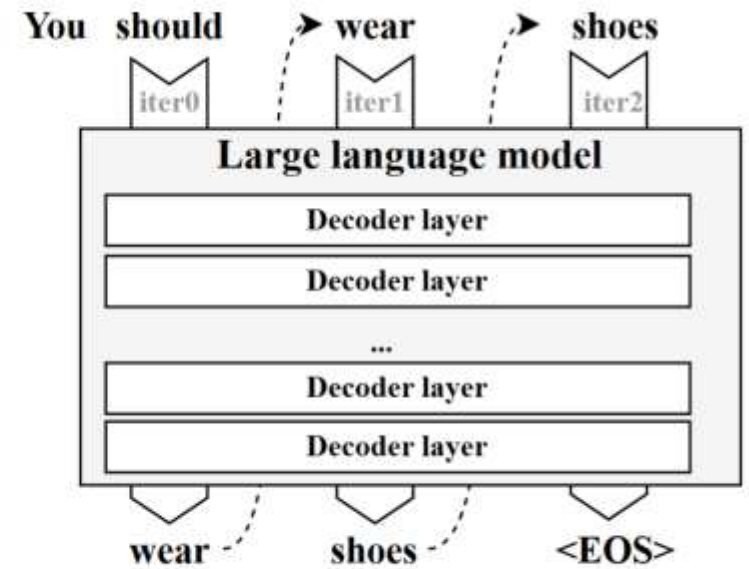
14 Dec 2023

# Background

- The generation of one token depends on previous tokens in autoregressive inference
- Some steps are easy and the other steps are hard



(a) GPT3 architecture




(b) Autoregressive inference

# Background

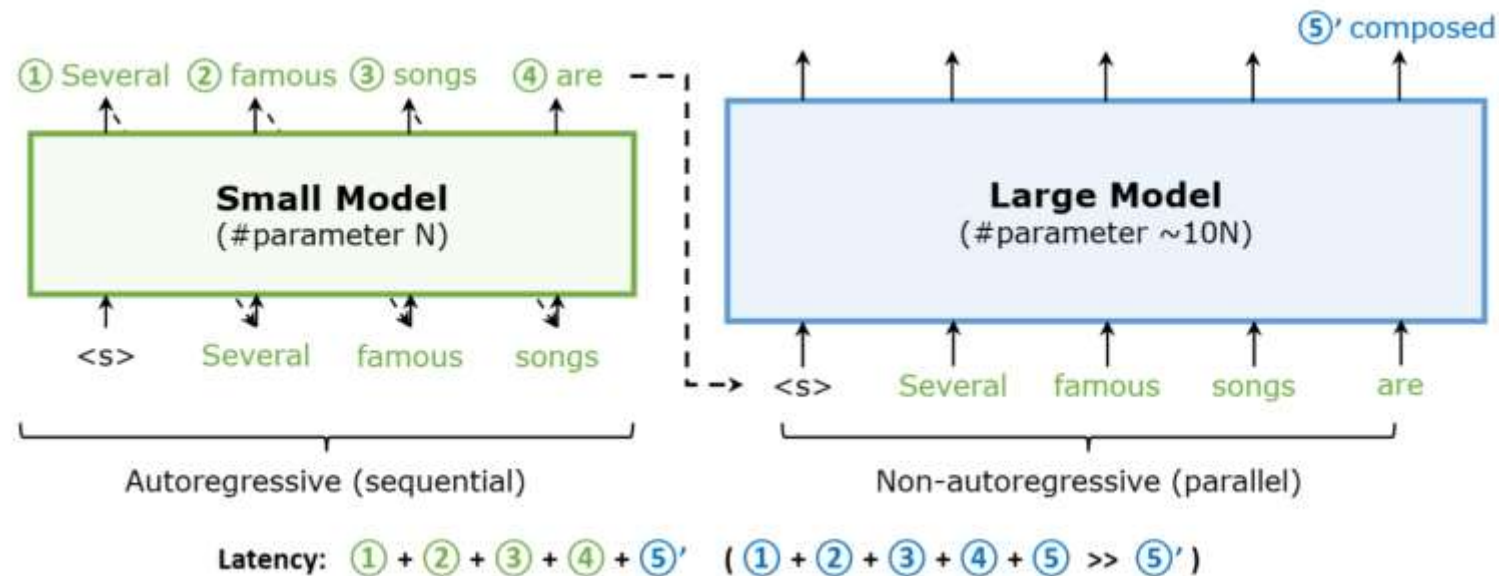
- A LLM is usually released with different size models
- The inference of a small model can be much faster than that of a large model

Model Name	$n_{\text{params}}$	$n_{\text{layers}}$	$d_{\text{model}}$	$n_{\text{heads}}$	$d_{\text{head}}$
GPT-3 Small	125M	12	768	12	64
GPT-3 Medium	350M	24	1024	16	64
GPT-3 Large	760M	24	1536	16	96
GPT-3 XL	1.3B	24	2048	24	128
GPT-3 2.7B	2.7B	32	2560	32	80
GPT-3 6.7B	6.7B	32	4096	32	128
GPT-3 13B	13.0B	40	5140	40	128
GPT-3 175B or “GPT-3”	175.0B	96	12288	96	128



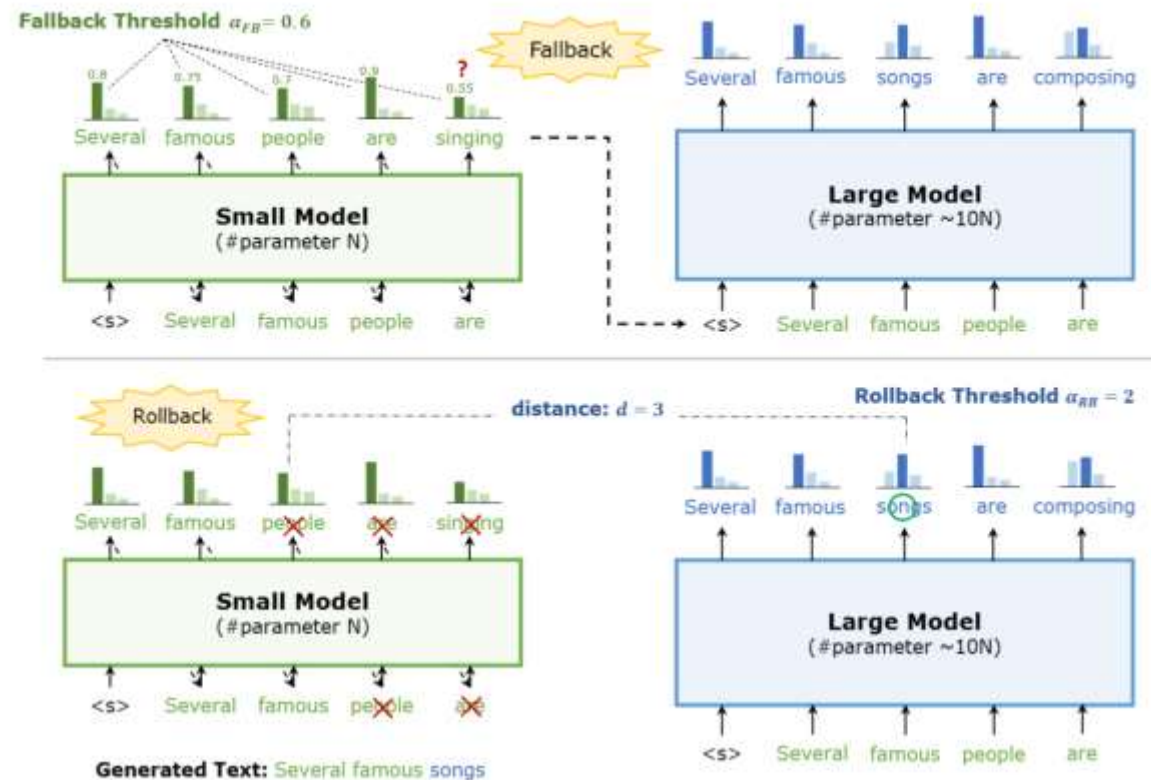
# Speculative Decoding

- Firstly, let a small model(draft model) sample tokens
- Then let the target model verify sampled tokens in parallel



# Speculative Decoding: Roll back

- If not all tokens are verified successfully, roll back to the last succeed token and continue to generate tokens.



# Speculative Decoding: Algorithm

- $\gamma$ : number of sampled tokens
- $r_i$  is sampled from  $U(0,1)$
- $M_p$ : Target model,  $M_q$ : Draft model
- A token is accepted if  $r_i < \max(1, \frac{p(x)}{q(x)})$
- In one Iteration, the draft model will generate  $\gamma$  tokens and the target model will generate 1 token.
- The final generated token has the same distribution as the target model

---

**Algorithm 1** SpeculativeDecodingStep

---

**Inputs:**  $M_p, M_q, prefix$ .

▷ **Sample  $\gamma$  guesses  $x_1, \dots, x_\gamma$  from  $M_q$  autoregressively.**

**for  $i = 1$  to  $\gamma$  do**

$q_i(x) \leftarrow M_q(prefix + [x_1, \dots, x_{i-1}])$

$x_i \sim q_i(x)$

**end for**

▷ **Run  $M_p$  in parallel.**

$p_1(x), \dots, p_{\gamma+1}(x) \leftarrow$

$M_p(prefix), \dots, M_p(prefix + [x_1, \dots, x_\gamma])$

▷ **Determine the number of accepted guesses  $n$ .**

$r_1 \sim U(0, 1), \dots, r_\gamma \sim U(0, 1)$

$n \leftarrow \min(\{i - 1 \mid 1 \leq i \leq \gamma, r_i > \frac{p_i(x)}{q_i(x)}\} \cup \{\gamma\})$

▷ **Adjust the distribution from  $M_p$  if needed.**

$p'(x) \leftarrow p_{n+1}(x)$

**if  $n < \gamma$  then**

$p'(x) \leftarrow \text{norm}(\max(0, p_{n+1}(x) - q_{n+1}(x)))$

**end if**

▷ **Return one token from  $M_p$ , and  $n$  tokens from  $M_q$ .**


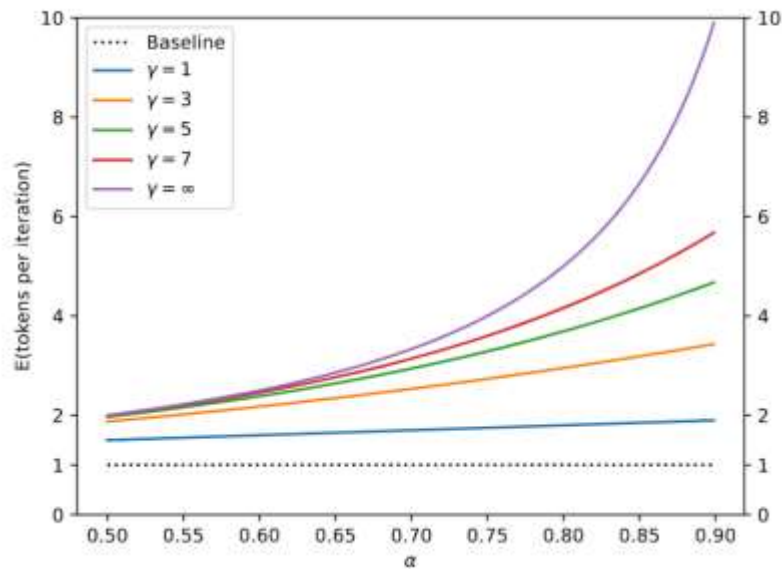
$t \sim p'(x)$

**return**  $prefix + [x_1, \dots, x_n, t]$

---

# Speculative Decoding: Analysis

- $\alpha$  is the average acceptance rate
- $c$  is the ratio between the time for a single run of the draft model and the time for a single run of target model



# Speculative Decoding

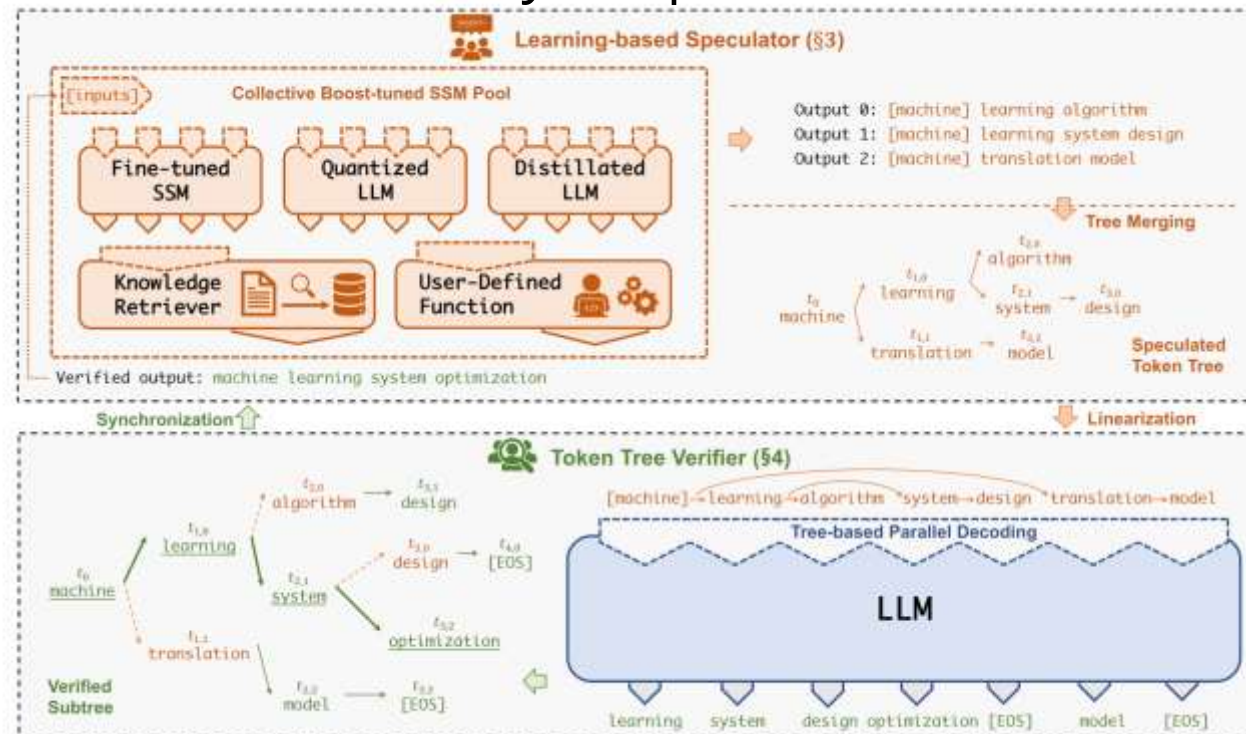
- Evaluation
- Target model:T5-XXL 11B

TASK	$M_q$	TEMP	$\gamma$	$\alpha$	SPEED
ENDE	T5-SMALL ★	0	7	0.75	<b>3.4X</b>
ENDE	T5-BASE	0	7	0.8	2.8X
ENDE	T5-LARGE	0	7	0.82	1.7X
ENDE	T5-SMALL ★	1	7	0.62	<b>2.6X</b>
ENDE	T5-BASE	1	5	0.68	2.4X
ENDE	T5-LARGE	1	3	0.71	1.4X
CNNDM	T5-SMALL ★	0	5	0.65	<b>3.1X</b>
CNNDM	T5-BASE	0	5	0.73	3.0X
CNNDM	T5-LARGE	0	3	0.74	2.2X
CNNDM	T5-SMALL ★	1	5	0.53	<b>2.3X</b>
CNNDM	T5-BASE	1	3	0.55	2.2X
CNNDM	T5-LARGE	1	3	0.56	1.7X



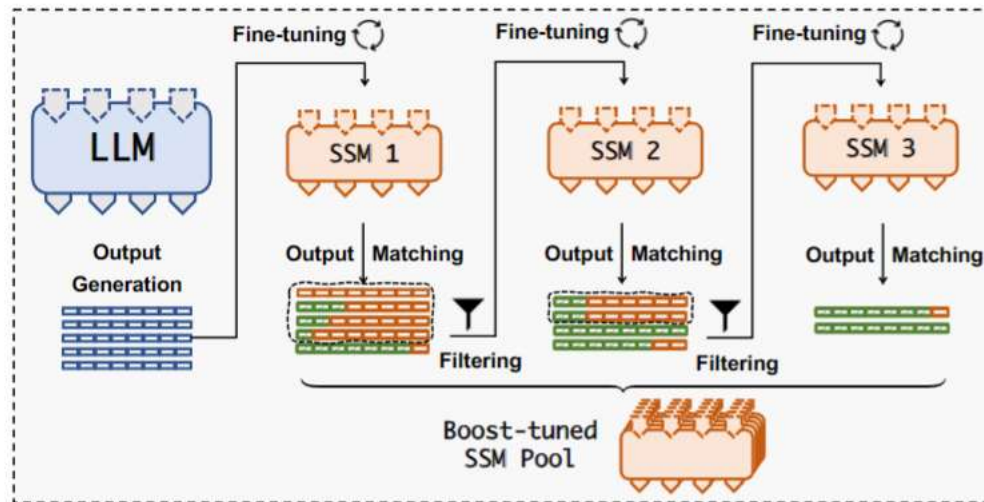
# SpecInfer

- Combine various collectively boost-tuned small draft models to jointly predict the LLM's output
- Use the target model to verify all predictions




# SpecInfer : Decoding

- Collectively boost-tune the draft models
- When decoding, every draft model generate a sequence
- The sequences will be merged into a tree




1. machine learning algorithm  
2. machine learning system design  
3. machine translation model

Token sequences



Token tree

- A topology-aware casual mask based on the token tree is used in self-attention.
- Target model will take all speculated tokens as input



# SpecInfer: Algorithm


- VerifyGreedy: Always choose the token with maximum probability
- VerifyStochastic: Sample tokens from a probability distribution

```
1: Input: A sequence of input tokens  $\mathcal{I}$ 
2: Output: A sequence of generated tokens
3:  $\mathcal{S} = \mathcal{I}$ 
4: while true do
5:    $\mathcal{N} = \text{SPECULATE}(\mathcal{S})$ 
6:    $\mathcal{O} = \text{TREEPARALLELDECODE}(\text{LLM}, \mathcal{N})$ 
7:   if use greedy decoding then
8:      $\mathcal{V} = \text{VERIFYGREEDY}(\mathcal{O}, \mathcal{N})$ 
9:   else
10:     $\mathcal{V} = \text{VERIFYSTOCHASTIC}(\mathcal{O}, \mathcal{N})$ 
11:   for  $t \in \mathcal{V}$  do
12:      $\mathcal{S}.\text{append}(t)$ 
13:     if  $t = \langle \text{EOS} \rangle$  then
14:       return  $\mathcal{S}$ 
15:
16: function VERIFYGREEDY( $\mathcal{O}, \mathcal{N}$ )
17:    $\mathcal{V} = \emptyset, u \leftarrow$  the root of token tree  $\mathcal{N}$ 
18:   while  $\exists v \in \mathcal{N}. p_v = u$  and  $t_v = \mathcal{O}(u)$  do
19:      $u = v$ 
20:      $\mathcal{V}.\text{append}(t_v)$ 
21:    $\mathcal{V}.\text{append}(\mathcal{O}(u))$ 
22:   return  $\mathcal{V}$ 
23:
```

```
24: function VERIFYSTOCHASTIC( $\mathcal{O}, \mathcal{N}$ )
25:    $\mathcal{V} = \emptyset, u \leftarrow$  the root of token tree  $\mathcal{N}$ 
26:   while  $u$  is a non-leaf node do
27:      $\mathcal{H} = \text{child}(u)$   $\triangleright$  The set of child nodes for  $u$ 
28:     while  $\mathcal{H}$  is not empty do
29:        $s \sim \text{rand}(\mathcal{H}), r \sim U(0, 1), x_s = \mathcal{H}[s]$ 
30:       if  $r \leq P(x_s | u, \Theta_{\text{LLM}}) / P(x_s | u, \Theta_{\text{SSM}_s})$  then
31:          $\triangleright$  Token  $x_s$  passes verification.
32:          $\mathcal{V}.\text{append}(x_s)$ 
33:          $u = s$ 
34:         break
35:       else
36:          $\triangleright$  Normalize the residual  $P(x | u, \Theta_{\text{LLM}})$ 
37:          $P(x | u, \Theta_{\text{LLM}}) := \text{norm}(\max(0, P(x |$ 
38:            $u, \Theta_{\text{LLM}}) - P(x | u, \Theta_{\text{SSM}_s})))$ 
39:          $\mathcal{H}.\text{pop}(s)$ 
40:       if  $\mathcal{H}$  is empty then
41:         break
42:        $\triangleright$  All SSMs fail verification; sample the next token
43:        $x_{\text{next}} \sim P(x | u, \Theta_{\text{LLM}})$ 
44:        $\mathcal{V}.\text{append}(x_{\text{next}})$ 
45:   return  $\mathcal{V}$ 
```

# SpecInfer: Evaluation

- 2 draft models, speculation length=16
- Verify mode is not mentioned

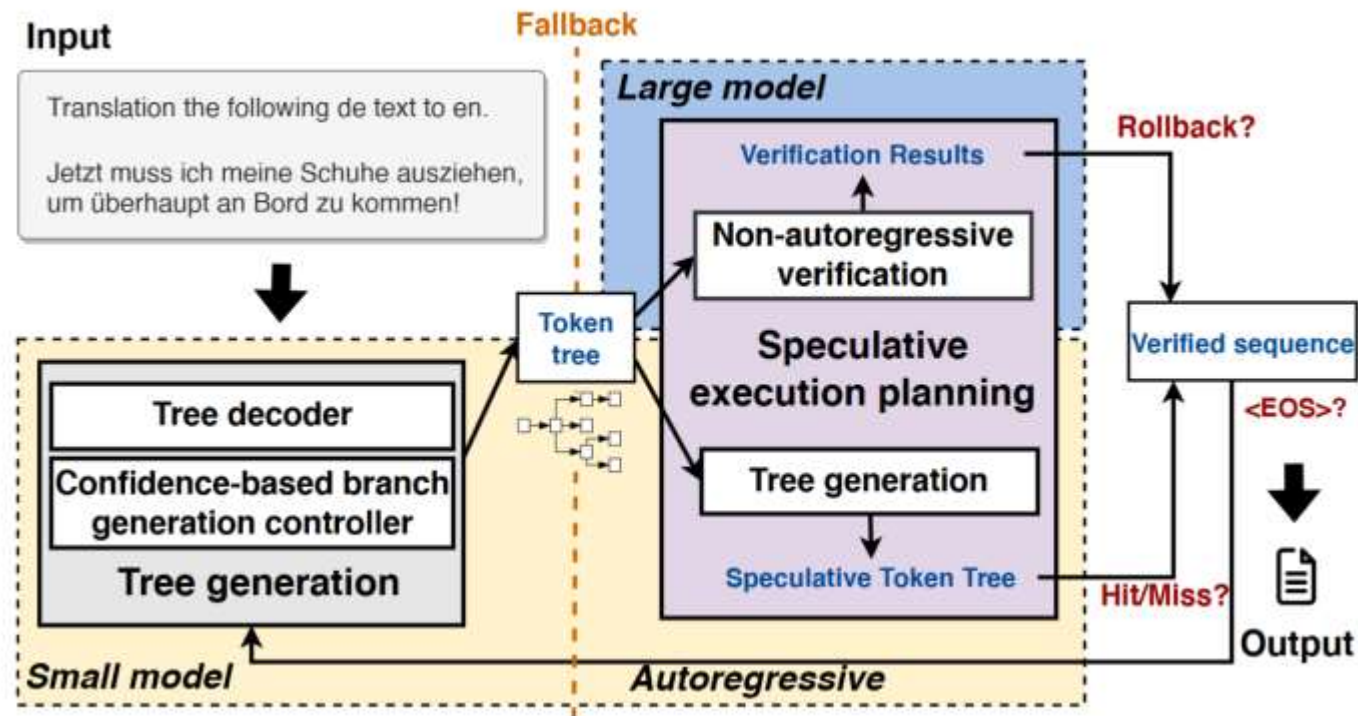


# SSMs	1	2	3	4	5
CIP	3.35	3.74	3.97	4.05	<b>4.11</b>
CP	2.71	3.14	3.32	3.45	<b>3.51</b>
WebQA	2.84	3.08	3.20	3.27	<b>3.31</b>
Alpaca	2.70	3.19	3.36	3.44	<b>3.49</b>
PIQA	2.98	3.21	3.36	3.44	<b>3.49</b>
Avg	2.92	3.27	3.44	3.53	<b>3.58</b>




# LLMCad

- On-device LLM inference
- Use one draft model to generate token tree



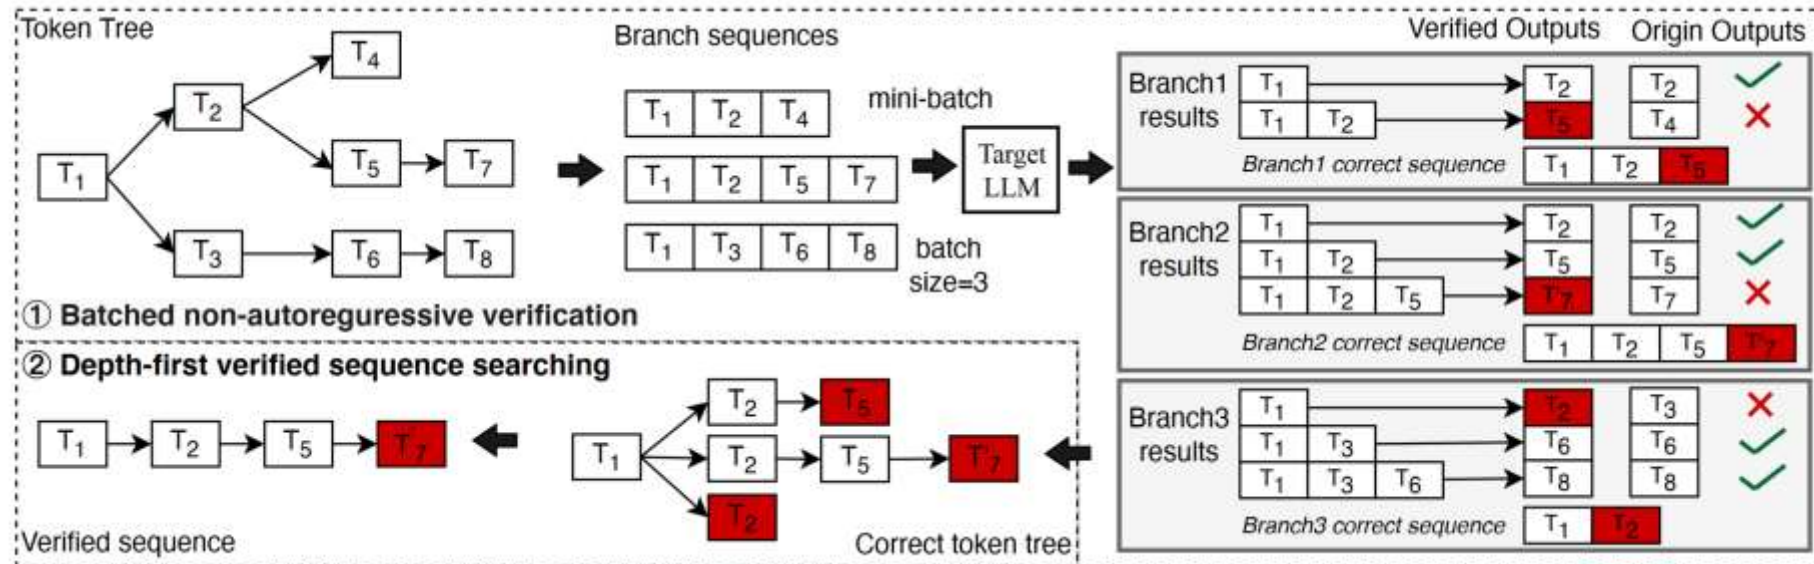
# LLMCad: Decoding

- Only one draft model to sample tokens
- Decide which branch to continue by  $f(x) = M * \frac{C_x}{\sum_{i=0}^N C_i} - T_x^B$
- Get the branch's context by a mask



# LLMCad: Verification

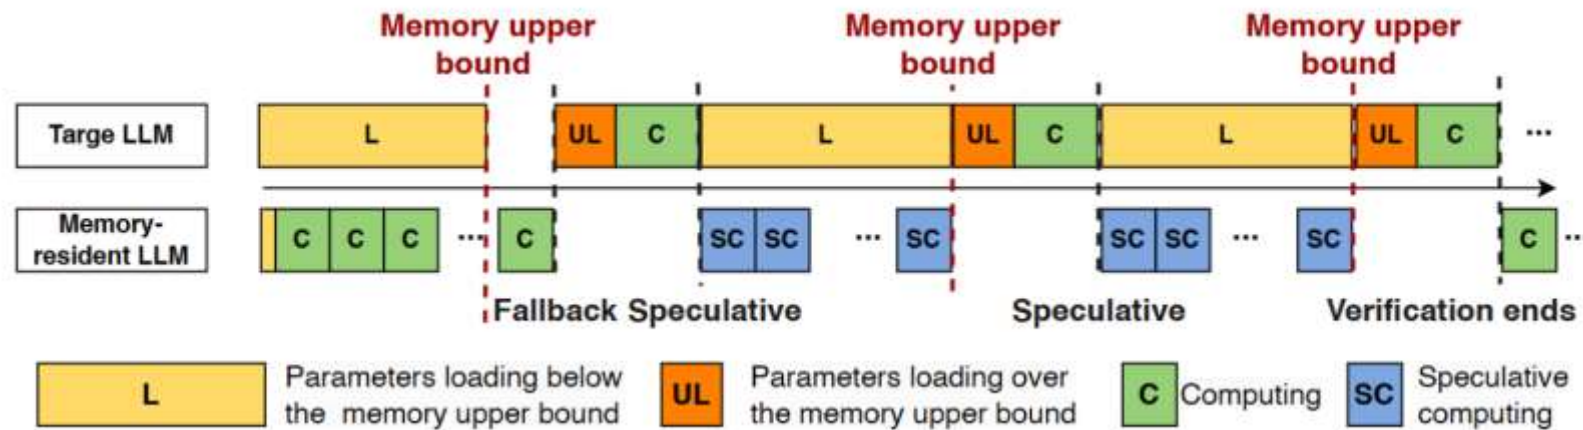
- self-adaptive threshold  $\alpha_{i+1} = \begin{cases} \alpha_i * 0.5 & \text{if } N_{correct} == N_{all} \\ \alpha_i / T_c \frac{N_{all} - N_{correct}}{N_{all}} & \text{if } N_{correct} < N_{all} \end{cases}$
- Verify when  $T_c = \max_{i=1}^{N_c} C_i < \alpha$
- Not tree mask but mini-batch





# LLMCad: Speculative Generation Pipeline

- Allow draft model to decode in the verification process
- To avoid two LLMs memory contention, parameters loading will stop before the memory upper bound is exceeded
- Speculative execution only when below the memory budget



# Other work

- Multi-stage draft models
- Medusa: Blockwise decoding + Token-tree
- Lookahead Decoding: Jacobi decoding + n-grams
- ...

# Conclusion

- Most of current work in this direction is still focusing on algorithm
- The speculative decoding algorithm itself is not yet fixed
- Need to do some profile to find some points of improvement

# Reference

- Fast Inference from Transformers via Speculative Decoding
- SpecInfer: Accelerating Generative Large Language Model Serving with Speculative Inference and Token Tree Verification
- LLMCad: Fast and Scalable On-device Large Language Model Inference